Human-Computer Interaction Lab / University of Maryland

home research publications academics about hcil members partnerships contact

Dynamic queries, starfield displays, and the path to Spotfire

Ben Shneiderman, February 4, 1999

The old days of command line interfaces and submitting queries to databases are passing quickly. In their place are dynamic queries and starfield displays that update a two-dimensional graphical display in 100 milliseconds. As users adjust sliders, buttons, check boxes, and other control widgets the starfield display containing color- and size-coded points us updated rapidly. Users feel they are in control and there is no more "RUN" button.

Two early applications of dynamic queries were built at the University of Maryland's Human-Computer Interaction Laboratory (HCIL Video Reports 1992) -- a chemical table of elements (HCIL TR 91-11) and a real estate HomeFinder (HCIL TR 92-01) (free downloadable version at http://www.cs.umd.edu/hcil/pubs/products.shtml).

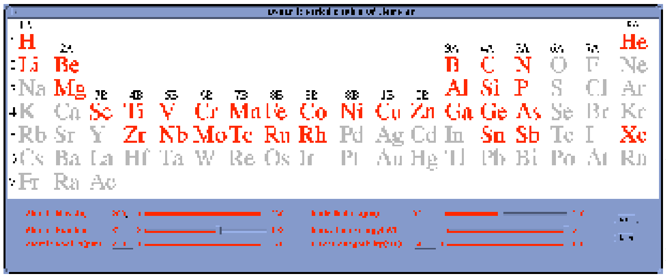
H He
Li Be B C N O F Ne
Na Mg Al Si P S Cl Ar
K Ca Sc Ti V Cr Mn Fe Co Ni Cu Zn Ga Ge As Se Br Kr
Rb Sr Y Zr Nb Mo Tc Ru Rh Pd Ag Cd In Sn Sb Te I Xe
Cs Ba La Hf Ta W Re Os Ir Pt Au Hg Tl Pb Bi Po At Rn
Fr Ra Ac

Christopher Ahlberg, a visiting student from Sweden during summer 1991, took up my lunch-time challenge to work on dynamic interfaces that applied direct manipulation principles. His first overnight success was making a modern slider-based version of a polynomial viewer that I had built as a graduate student in 1972 (The Mathematics Teacher 67,2 (February 1974), pages 111-113). As users move the sliders for each coefficient, the curve gracefully reshapes on the screen - dancing parabolas. Within a week he had satisfied my second challenge of a dynamic query for the chemical table of elements. He put up the periodic table with chemical symbols in red and six sliders for attributes such as atomic radius, ionization energy, and electronegativity. As users move the sliders, the chemical symbols change to red showing the clusters, jumps, and gaps that chemists find fascinating. A study with 18 chemistry students showed faster performance with use of a visual display (versus a simple textual list) and a visual input device (versus a form fillin box).

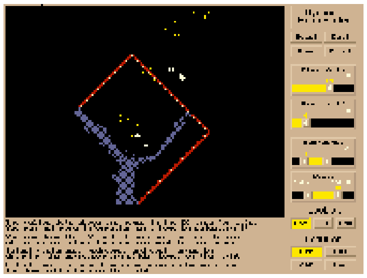

Christopher Williamson's HomeFinder showed a map of Washington, DC and 1100 points of light indicating homes for sale. Users could mark the workplace for both members of a couple and then adjust sliders to select circular areas of varying radii. Other sliders selected number of bedrooms and cost, with buttons for air conditioning, garage, etc. Within seconds users could see how many homes matched their query, and adjust accordingly. Controlled experiments with benchmark tasks showed dramatic speed-ups in performance and high subjective satisfaction (1, 2, 3) (HCIL TR 93-01and HCIL TR 94-16) (HCIL Video Reports 1993 and 1994). This demo continues to be one of the most compelling and comprehensible even though it is almost 8 years old.

Williamson earned a trip to the ACM SIGIR'92 conference in Copenhagen to present his work. Then he went on to the University of Colorado at Boulder to do a master's thesis that expanded the idea into a well-engineered and commercial viable version. One of the amusing stories about this project was the unwillingness of corporate or university sources for regional housing information to share their data. Undaunted, Chris Williamson and his friends took a Sunday Washington Post and typed in the data for the 1100 homes. One of the amazing stories is the resistance of these same institutions to learn about or apply our approach.

The concept of a generic two-dimensional scattergram with zooming, color coding, and filtering was first applied in the FilmFinder (4) (HCIL TR 93-14) (HCIL Video Reports 1994 and repeated in 1996). Our lab was working on interactive TV applications and we had a brainstorming session in the conference room with about eight attendees including Chris Ahlberg, who joined us for a second summer. I asked each person to describe a possible interface for finding a film from a library of 10,000 videotapes. As the variants of traditional approaches with command lines and menus were rejected, it became more difficult for each speaker to come up with something fresh. Since every alternative was text based, I concluded the session by proposing a two dimensional layout with years on the x-axis and popularity on the y-axis. The idea was quickly accepted and refined.

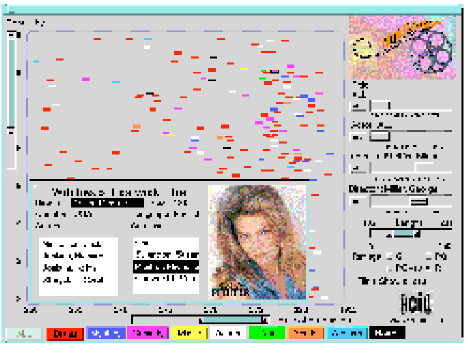

By the next morning Chris had a prototype showing fifteen hundred films with color coded spots (red for drama, white for action, etc.). As the weeks passed and other students built components, Chris integrated them into the FilmFinder. A range slider allowed filtering by the length of the film and buttons allowed selection by ratings. A click on one of the spots produced a pop-up box with details of each film and a picture of one of the actors/actresses. The picture for about 50 actors were grabbed from the net, and Michelle Pfeiffer (Chris's choice) became the default for the others. Demos were carefully scripted around the pictures that we had.

The idea for the alphaslider had been germinating in our lab for two years, but we were stuck with designs that had one item per pixel (HCIL TR 93-08), limiting its use to a few hundred items. Chris proposed the concept that many items could be tied to a single pixel, but I thought he misunderstood the technology - it was me who just couldn't grasp his solution. It took me a few minutes to shake free from my assumptions and then it was immediately clear that his approach would work. Chris built four versions and ran a study with 24 subjects (5) (HCIL TR 93-15). The FilmFinder used the alphasliders to select form thousands of actors, actresses, and directors. The alphaslider idea spawned a variety of improvements and further empirical studies.

The FilmFinder videotape was also produced during Chris's intense 12-week summer visit and it remains my most popular tape to show, even 5-6 years later. The group became known as the Widget Carvers of College Park, commemorated by a hand carved wooden sign that hangs at our lab's doorway. Later work on the FilmFinder sought to improve the data structures (HCIL TR 93-16) and make the zooming smoother (HCIL TR 93-06). Our advanced research continues to increase the size of the database while keeping response times rapid (8, 9).

Chris Ahlberg returned to Sweden to work on his excellent PhD on Dynamic Queries at Chalmers University. He developed an enhanced UNIX implementation, called IVEE (Information Visualization and Exploration Environment) that was flexible in reassigning the axes to other attributes and had a scroll bar to permit large numbers of sliders. It allowed importation of arbitrary flat files (6) and therefore was applied in many projects, including the State of Maryland's Department of Juvenile Justice (7) (HCIL TR 96-15)(HCIL Video Reports 1995) .

Ahlberg gathered his friends, and found venture capital to start a company. The commercial version of the starfield display, now called Spotfire, allows increased user control and greater flexibility. Spotfire was launched in mid-1996 by IVEE Development, which became renamed as Spotfire Inc. (http://www.spotfire.com). It maintains a development center in Goteborg, Sweden with its North American headquarters in Cambridge, MA. Because of my sense of paternity for the idea, I agreed to be on the board and have enjoyed watching Chris Ahlberg's transformation into a successful businessperson. Spotfire has become a leader in visual data mining and information visualization and is available in Windows and Windows NT versions. It has become enormously successful in the challenging pharmaceutical drug discovery task, leading 16 of the 20 big

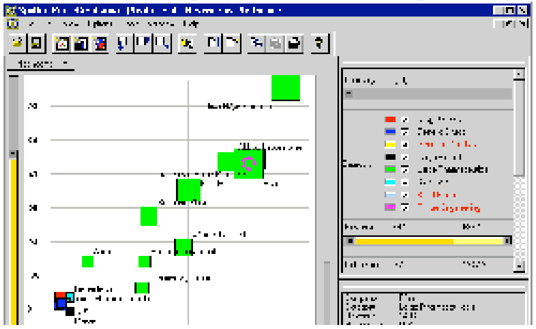

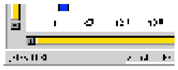

pharmaceutical companies to become major adopters. Easy import/export of data, rapid change to axes, color coding, or size coding, and collaboration support have made Spotfire a leader in its class. Expansion to other application domains is proceeding.

The concept of dynamic queries became the basis for applications such as pruning of large tree structures based on sliders tied to attributes of nodes (HCIL 95-12) (HCIL Video Reports 1995) and a Youth services database for the Maryland Dept. of Juvenile Justice (HCIL TR 96-07) (HCIL Video Reports 1995). Our recent efforts have applied Spotfire to web site log visualization, with the goal of understanding patterns of usage for ecommerce.

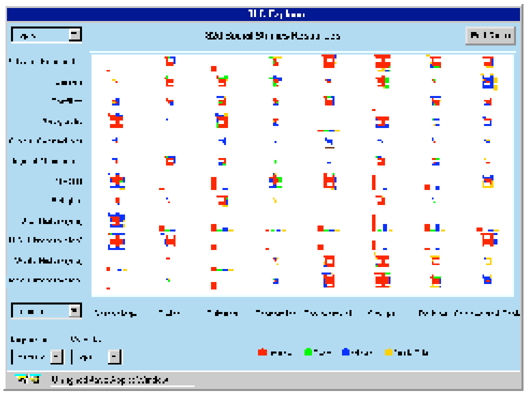

In developing a Java-based visual tool for high school teachers to find educational resources, the Baltimore Learning Community (www.learn.umd.edu) we found that a simplified starfield display that had a discrete grid, rather than continuous axes, was easier to use (HCIL TR 97-15). In this database, color coding would typically be used for media type, for example, red for videos, green for web sites, blue for images, and yellow for texts. When more than 49 documents wind up in a grid box, Dotfire shifts to a color-coded bar graph showing the frequency counts for each type of document. The typical axes for this application are topics (Arts and Humanities, Careers, Conflicts, Geography, Religion, U.S. History, World History) and outcomes (Chronlogy, Civis, Cultures, Economics, Environment, Politics, Science and Technology). Other categorical axes used in this application were year (1991 to 1998) and source (Discovery Communications, National Geographic, Library of Congress, etc.). Users could control and quickly change the axes and the color coding.

Another extension of dynamic queries is to query previews for NASA remote sensing environmental databases. Instead of showing results within 100 milliseconds, the query preview simply shows the size of the result set. This eliminates wasted time with empty and with overly large result sets. This approach works well on the World Wide Web and it has been implemented on NASA websites and elsewhere.

REFERENCES

1) Ahlberg, Christopher, Williamson, Christopher, and Shneiderman, Ben, Dynamic queries for information exploration: An implementation and evaluation, Proc. ACM CHI'92: Human Factors in Computing Systems (1992), 619-626. (link to HCIL TR 92-01)

2) Williamson, Christopher, and Shneiderman, Ben, 1992. The Dynamic HomeFinder: Evaluating dynamic queries in a real-estate information exploration system, Proc. ACM SIGIR'92 Conference, Copenhagen, Denmark, (June 1992), 338-346. Reprinted in Shneiderman, B. (Editor), Sparks of Innovation in Human-Computer Interaction, Ablex Publishers, Norwood, NJ, (1993), 295-307.

3) Shneiderman, Ben, Dynamic queries for visual information seeking, IEEE Software 11, 6 (1994), 70-77.

4) Ahlberg, Christopher and Shneiderman, Ben, Visual Information Seeking: Tight coupling of dynamic query filters with starfield displays, Proc. of ACM CHI94 Conference (April 1994), 313-317 + color plates. Reprinted in Baecker, R. M., Grudin, J., Buxton, W. A. S., and Greenberg, S. (Editors), Readings in Human-Computer Interaction: Toward the Year 2000, Second Edition, Morgan Kaufmann Publishers, Inc., San Francisco, CA (1995), 450-456.

5) Ahlberg, Christopher and Shneiderman, Ben, AlphaSlider: A compact and rapid selector, Proc. of ACM CHI94 Conference, (April 1994), 365-371.

6) Ahlberg, Christopher and Wistrand, Erik, IVEE: An information visualization & exploration environment, Proc. IEEE Information Visualization '95, IEEE Computer Press, Los Alamitos, CA (1995), 66-73.

7) Ellis, Jason, Rose, Anne, Plaisant, Catherine, Putting visualization to work: ProgramFinder for youth placement, ACM CHI 97: Human Factors in Computing Systems , ACM, New York (March 1997), 502-509

8) Tanin, Egemen, Beigel, Richard, and Shneiderman, Ben, Design and evaluation of incremental data structures and algorithms for dynamic query interfaces, IEEE Information Visualization Conference (October 1997).

9) Tanin, Egemen, Beigel, Richard, and Shneiderman, Ben, Incremental data structures and algorithms for dynamic query interfaces, ACM SIGMOD Record 25, 4 (December 1996), 21-24.

Refereed Videos available from ACM SIGCHI (http://www.acm.org/sigchi)

Ahlberg, Christopher , Shneiderman, Ben, and Williamson, Christopher , Dynamic Queries, ACM SIGGRAPH Video Review 77, 10 min. (1991).

Ahlberg, Christopher and Shneiderman, Ben, Visual Information Seeking: Tight coupling of dynamic query filters with starfield displays, ACM SIGGRAPH Video Review 97, 7 min.(April 1994).

ELSEVIER

Interacting with Computers 11 (1998) 137–146

Interacting with Computers

The Attribute Explorer: information synthesis via exploration

Robert Spence*, Lisa Tweedie
Imperial College, London, UK

Accepted 12 February 1998

Abstract

The Attribute Explorer is a visualization tool in which the graphical and interactive presentation of data supports the human acquisition of insight into that data. The underlying concept employed is that of interactive linked histograms. The advantage of the Attribute Explorer derives from its ability to support both qualitative exploration and quantitative design decisions, as well as a smooth transition between these two activities. © 1998 Published by Elsevier Science B.V. All rights reserved

Keywords: Attribute Explorer; Qualitative exploration; Quantitative decision making

1. Data and information

1.1. The traditional approach

The conventional approach to so-called ‘information retrieval’ appears to be based upon a number of tacit assumptions. It is assumed, for example, that the user is generally familiar with the relevant database(s) and therefore has no need to gain such familiarity. It is also assumed that the user will come to the database with a precisely formulated question, and that the answer will usually provide a satisfactory solution. There will therefore be little need to formulate a revised question. Furthermore, it is usually assumed that the user has been trained to access the database.

1.2. Information synthesis

Like many others, we take the view that the traditional approach to ‘information retrieval’ is relevant only to an extremely small fraction of ‘real world’ problems. In most cases, the user—whether novice or professional—may not be familiar with the

* Corresponding author.

PII: S0953-5438(98)00022-8

database, and will need to explore its contents to gain that overall insight needed to formulate a useful question (‘‘X seems to trade-off against Y’’; ‘‘Z seems to be bimodally related to W except for…’’). Having gained such an overview, the user will in all probability revise the question. Indeed, we suggest that problem formulation is a major activity and takes place concurrently with, and is an indispensable component of, problem solution. Indeed, a principal objective of visualization tools such as the Attribute Explorer is to inform the user in such a way that the problem formulation can be refined. We therefore prefer to speak of ‘information synthesis’ rather than information retrieval, since the latter term might be assumed to imply that information can be instantly retrieved, whereas we believe that, through exploration, the user is gradually informed. A similar philosophy appears to be implicit in the work of many others, e.g. [1–3].

1.3. Exploration

Our thesis can be illustrated by the task of home finding. The prospective buyer will enter an estate agent’s office with only an imprecisely formulated need (‘‘not more than £60K, hopefully more than two bedrooms and quite close to a railway station’’) and with little or no knowledge of available houses. The first need here is to explore the available houses to gain insight into general characteristics (‘‘prices increase steeply near that park’’; ‘‘there’s no chance of a third bedroom at £60K’’), so that a more focussed question can be posed. Exploration of this kind will eventually lead to the identification of a house deemed acceptable, though it may not satisfy many or even any of the original subjectively expressed requirements. Similar remarks apply to a wide variety of tasks, for example the problem faced by a marketing manager within an investment house in deciding upon a new marketing strategy on the basis of accumulated data.

1.4. The Attribute Explorer

Our consideration of information synthesis has led to the invention of a new tool called the Attribute Explorer. The task to which it is directed, and which has been described by the example above, can be stated more formally as

> Given a collection of **objects**, each described by the values associated with a set of **attributes**, find the most acceptable such object or, perhaps, a small number of candidate objects suited to more detailed consideration.

As already pointed out, the objects may be as different as houses and marketing schemes. Similarly, the definition of ‘most acceptable’ is purposely vague because the refinement of a problem is an essential part of the process of choosing one object from many on the basis of its attributes. The Attribute Explorer [4] is so named because that refinement is achieved primarily by exploration.

1.5. Visualization and externalization

The Attribute Explorer is characterized by its graphical representation and presentation of data, and by the nature of the interaction it supports. It is the interactive and graphical exploration of data that, in our opinion, facilitates the user's acquisition of insight into that data. In this context, we are reminded of a Chinese proverb to which we have added the last line:

> I hear and I forget, I see and I remember, I interact and I understand, *I interact* ***responsively*** *and I discover*.

where, by 'responsively', we mean that an effect is perceived in less than about 0.1 s following its cause.

A key feature of the Attribute Explorer is that it forms an external representation of the user's problem. Constraints relevant to the problem are hard-coded into the tool's functionality so that the user does not need to focus on remembering those constraints. In this way, the rules of the problem are made implicit and the user can concentrate directly on the solution of their problem [5]. External representations of this kind are useful because they also serve as a constant reminder of the current state of the solution without recourse to memory.

To establish the salient aspects of the Attribute Explorer, we first briefly examine a traditional information retrieval scheme to identify its shortcomings. We then describe a tool (Dynamic Queries) which, in 1992, heralded a new approach to information synthesis and which stimulated the invention of the Attribute Explorer.

2. Conventional database queries

Continuing with the home-finding example, a first and major drawback associated with conventional 'information retrieval' interfaces such as those using SQL is that the user, who may well be a lay person, has to learn a language sufficiently well to be able, for example, to enter

select house address
from my database
where price $<\ =100\,000$ **and**
bathrooms $=2$ **and**
bedrooms $>\ =3$

to indicate an interest in finding a house costing no more than £100K with two bathrooms and no less than three bedrooms. A second drawback is that errors in formulating the query are usually not tolerated, with devastating consequences for the novice user. A third drawback is encountered if, as often happens, no items satisfying the query are found:

0 HITS.

Equally unhelpful is the situation in which too many items are identified:

10 573 HITS.

In both of these cases, but also in many other situations, a fourth drawback becomes apparent: no guidance is offered as to how the query may usefully be modified to obtain not only a manageable number of hits, but 'movement' towards a useful problem statement. A fifth, and very serious, drawback is that no data are presented unless specifically requested: data are being unnecessarily hidden, with a consequent loss of contextual information. A sixth drawback is that it is exceedingly difficult for the user to synthesize knowledge of the contents of the database, and to discover such global characteristics as trade-offs (e.g., between price and distance to station) and correlations (e.g., between price and number of bedrooms). Thus, six drawbacks associated with conventional database query interfaces have to be addressed:

- The discretionary user must learn a language.
- Errors are not tolerated.
- Two few or too many hits may be registered.
- No guidance is given regarding potentially beneficial modification to the query.
- Useful contextual data are often unnecessarily hidden.
- It is difficult for a user to build an internal model.

As already stated, most of these drawbacks can be traced to a fundamental—and, in the vast majority of situations, unrealistic—assumption that ‘a user will typically know precisely what single question needs to be answered’ and, additionally, ‘knows a great deal about the content of the database’. Our rejection of this assumption is fundamental to the developments we report below.

3. Dynamic queries

Perhaps the most significant recent change in ‘information retrieval’ attitudes was evidenced by the concept of Dynamic Queries, first reported by Williamson and Shneiderman [6]. An effective illustration (Fig. 1) is again provided by a lay person’s search for a home. To the right of a map are a number of scales with adjustable sliding limits, each associated with a particular attribute: ‘house price’, ‘number of bedrooms’ and ‘journey time’ to work are typical examples. Adjustment of the upper and/or lower acceptable limits to any of these attributes causes the corresponding selection to be made from all houses in the database, and the result of that selection immediately displayed by dots on the map. A significant advantage is that responsive interaction—by which we mean that an effect occurs within less than about 0.1 s of its cause—not only allows answers to be obtained rapidly but, most importantly, supports dynamic exploration, often called the ‘What if?’ activity. Even though no direct indication of how a query might usefully be modified is given, some indication can, to some extent, now be obtained by manual exploratory variation of the attribute limits.

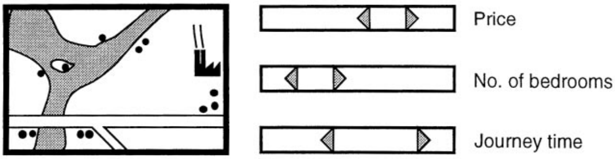

Fig. 1. The dynamic HomeFinder.

4. Disclosure of data

A major drawback to the dynamic queries concept as implemented by Williamson and Shneiderman [6] is that data are disclosed only when they satisfy the query. Especially when attention is—albeit temporarily—focussed on a small fraction of the available data, few data are available to provide **context** and **guidance** for later exploration. We therefore suggested that ‘all available data can and usually should be presented’. A convenient presentation could take the form of a histogram associated with each of the attribute scales. Since such a visualization tool is particularly valuable for exploring objects characterized by a number of attributes, the name Attribute Explorer was given to the tool that eventually emerged.

5. The Attribute Explorer

Fig. 2 shows the essential characteristics of the Attribute Explorer: the distribution of all objects (houses) on all attribute scales is represented by a histogram. Again, the example of house buying is adopted for illustration, and for simplicity only three attributes are considered. Attributes need not be continuous: they can be ordinal, categorical and/or discrete.

Each object is represented once—by a rectangle of suitable size—within each of the attribute histograms. A mouse-click on such a rectangle identifies the value of that object’s attributes on all the attribute scales, as shown in Fig. 2. An immediate advantage deriving from the display of ‘all’ the data is an awareness of the attribute ranges and the relative availability of objects with respect to attribute values: a further—and considerable—advantage, soon to be apparent, is the provision of context for specific queries.

5.1. Limits

With each attribute histogram is associated a scale containing two adjustable limits. These limits can be independently adjusted to specify the minimum and maximum acceptable values of an attribute. Alternatively, the bar linking these limits can be moved—taking the limits with it—to study the effect of moving a ‘range’ of attribute values. However the lower and upper limits are placed, the effect is the identification of all objects having attribute values between those limits. The objects so identified are then

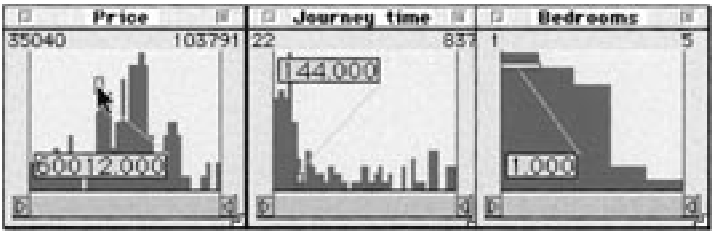

Fig. 2. The Attribute Explorer, showing the distribution over three attributes.

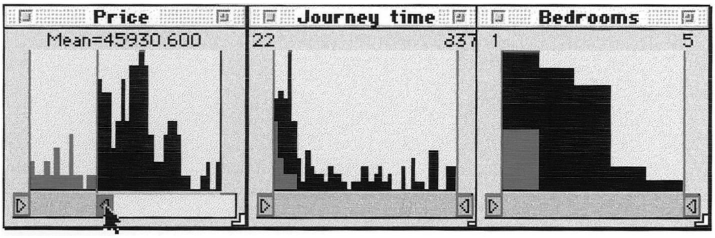

Fig. 3. Selection of a ‘Price’ range is reflected by colour coding into the remaining histograms.

marked by colour coding in the relevant section of the histogram, as shown for the ‘Price’ attribute in Fig. 3.

5.2. Attribute interaction

It is a simple matter, as illustrated by the ‘number of bedrooms’ and ‘journey time’ histograms in Fig. 3, to arrange that the same set of objects be colour-coded ‘on all other attribute histograms’, thereby providing an immensely valuable indication of the inter-relations between attributes. Manual variation, either of a limit or of a range defined by limits, leads to a corresponding alteration of the colour coding of the attribute histograms, additionally enhancing the opportunity to explore the inter-relation between attributes. Thus, manual adjustment of the ‘Price’ range from its lowest to its highest extent may well be seen to cause the ‘bunch’ of green houses in the ‘journey time’ histogram to ‘move’ in the opposite direction, suggesting a trade-off between these two attributes. To aid perception of such an effect, all colour-coded sections of the histogram are ‘lower justified’. The interaction between histograms is a special form of the ‘brushing’ technique first introduced by Newton [7] and later quite widely exploited [8].

5.3. Summary information

It is not always easy, during exploration, to interpret a ‘dynamic’ alteration of the colour of the separate elements of a histogram. To aid in such interpretation, summary information is available in the form of a yellow circle (Fig. 4) positioned on the scale at the average of the attribute values of the identified set of objects. Now, when the range of ‘Price’ (Fig. 3) is quickly moved up and down the corresponding scale, the yellow ‘average’

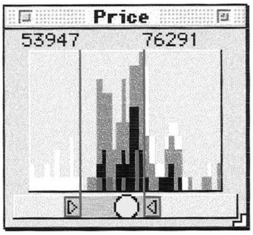

Fig. 4. The round dot indicates the average value for the selected houses.

circles on the ‘number of bedrooms’ and ‘journey time’ scales move correspondingly, allowing the relation between these attributes to be more readily discerned.

5.4. Multiple limits

As happens in the Dynamic Query display (Fig. 1), it is possible to arrange for a set of objects to be identified by limits placed on more than one attribute, and for the objects so identified to be colour-coded within each histogram, as shown in Fig. 5. Stated formally, the object selection shown is the result of performing the Boolean AND operation on the separate attribute ranges. Such a facility is particularly helpful when trying to identify one object—or a small number of candidate objects—to satisfy a particular need. The figure also demonstrates, by the inclusion of the map that would be expected of a house-finding tool, that attribute representations can be two-dimensional, in this case permitting the specification of a desirable geographical location.

5.5. Additive encoding: sensitivity information

A major disadvantage of both the Dynamic Queries interface (Fig. 1) and the example of an SQL interaction is the absence of any direct guidance as to how a ‘query’ may be modified to lead to more useful information, a need particularly acute in the ‘all-or-nothing’ situation of too many or too few hits. Such a disadvantage can be markedly reduced by additive encoding, in which colour coding is applied not only to those objects satisfying all attribute limits but, additionally and separately, to those failing one, two, three and more such limits. Thus, in the illustration of Fig. 5, and in addition to the colour-coded houses satisfying all limits, black indicates those that fail one limit, dark grey those that fail two limits, and so on. It is now apparent how a limit may usefully be changed: a limit extended to include a black object will cause it to turn green. Even for a query which identifies no acceptable (green) objects, Fig. 6 shows how limits might be relaxed to discover an acceptable house. Such information may be called ‘sensitivity information’, since colour indicates the reduction in the number of violated limits achieved by movement of a limit. This technique has much wider application than to the Attribute Explorer: it could be applied, for example, to Chalmers et al.’s [2] BEAD representation.

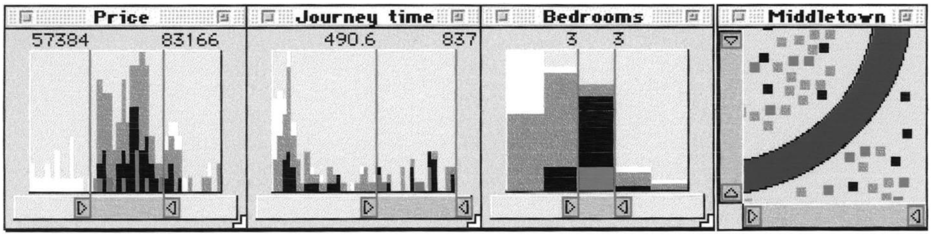

Fig. 5. Grey-scale coding shows how many limits are violated.

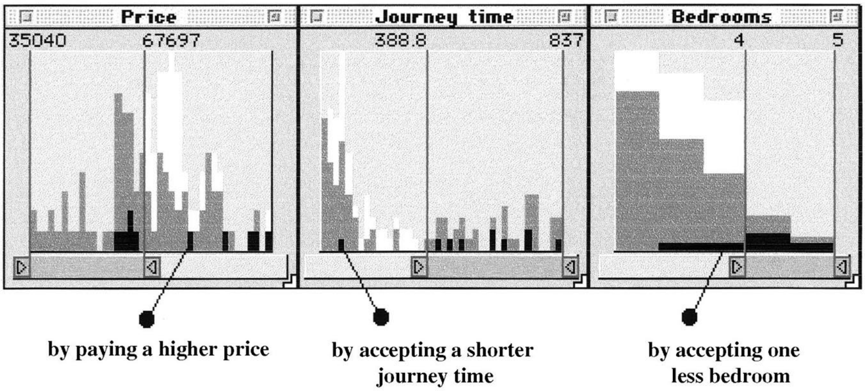

Fig. 6. In this case no houses satisfy the query, but sensitivity information indicates how limits could be altered to satisfy the query.

5.6. Previous work

Certain individual aspects of the Attribute Explorer are closely related to existing published work. The encoding of objects selected according to one attribute (e.g., price) in a display associated with another attribute (e.g., 'journey time', as in Fig. 3) is termed 'brushing' and was first introduced by Newton [7] in 1978. It is found in the Data Desk [9] visualization package. Eick [10], in 1994, reported his data visualization sliders combining a histogram with a slider mechanism, but did not explore the use of sliders as filtering mechanisms or for activating Boolean operations.

5.7. User control of link logic

The Attribute Explorer's facilitation of the Boolean AND function (leading, for example, to identification of the green houses in Fig. 5) prompts the question as to whether other Boolean operations on the attributes can conveniently be supported. Similarly, the valuable sensitivity information provided by colour coding could usefully be generalized. Both these potential enhancements to the Attribute Explorer can be provided via a single tool called the Link Crystal [11].

The Link Crystal was inspired by Spoerri's InfoCrystal [12] and, for a three-attribute situation, can be illustrated by derivation from a three-variable Venn diagram, as shown in

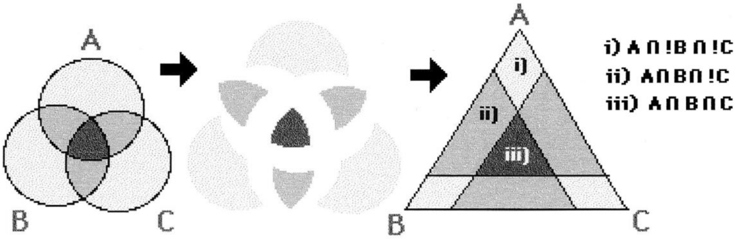

Fig. 7. Development of an InfoCrystal.

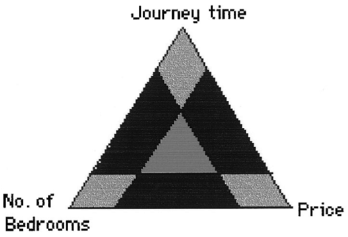

Fig. 8. A Boolean Link Crystal.

Fig. 7. Sub-regions of the Link Crystal correspond to different Boolean functions of the three attributes, and colour can be assigned to the sub-regions to control the colour coding of the displayed histograms. Thus, the colour coding shown in Fig. 5 corresponds to the Link Crystal shown in Fig. 8. If, on the other hand, interest is confined to houses that satisfy the price limits OR the limits on the number of bedrooms, the Link Crystal of Fig. 9 is relevant.

Selection of the colour associated with a sub-region of the Link Crystal is by mouse-click on that region, successive clicks cycling the colour through the available range. The concept of the Link Crystal, like that of the InfoCrystal, generalizes to any number of attributes. Whether or not a Link Crystal will be made available to the user will, of course, depend upon the nature of the user and the task for which the Attribute Explorer is intended.

6. Conclusions

Extensive discussions with a wide range of users have established that the features integrated within the Attribute Explorer support two major concurrent activities. One is **qualitative exploration** of the kind that allows insight to be gained into a body of data and which, for example, supports the discovery of new knowledge such as trade-offs and correlations. The other activity is that of **quantitative decision making** in which, as more detailed insight is acquired, small and usually iterative adjustments are made to attribute limits in order to identify an 'optimum' object or, perhaps, a small set of candidate objects worthy of more detailed study. Such a blend of qualitative exploration and quantitative decision making, especially with the smooth transition that can occur between

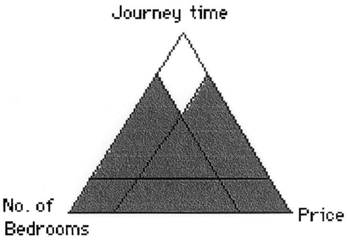

Fig. 9. A Link Crystal showing 'price OR number of bedrooms'.

these two activities in the Attribute Explorer, is conducive to the activity of information synthesis.

Apart from its generally favourable reception, no formal evaluation of the Attribute Explorer has been performed. Such an evaluation may be useful when systems employing this and other concepts have matured sufficiently.

The concept of linked interactive data histograms has also been applied to the field of design in the form of the Influence Explorer [13].

Acknowledgements

We gratefully acknowledge the valuable collaboration of John Nelder, Huw Dawkes and Zahid Malik. Lisa Tweedie was supported in her research by an EPSRC Research Studentship.

References

[1] C. Ahlberg, C. Williamson, B. Shneiderman, Dynamic queries for information exploration: an implementation and evaluation, ACM, Proceedings CHI'92, 1992, pp. 619–626.

[2] M. Chalmers, R. Ingram, C. Pfranger, Adding imageability feature to information displays, ACM, Proceedings UIST'96, 1996.

[3] M. Hearst, Tile bars: visualization of term distribution information in full text information access, ACM, Proceedings CHI'95, 1995, pp. 59–66.

[4] L. Tweedie, R. Spence, D. Williams, R. Bhogal, The Attribute Explorer, ACM, Video Proceedings CHI'94, 1994.

[5] J. Zhang, D.A. Norman, Representations in distributed cognitive tasks, Cognitive Science 18, 1994, 87–122.

[6] C. Williamson, B. Shneiderman, The Dynamic Homefinder: evaluating dynamic queries in a real estate information exploration system, ACM, Proceedings SIGIR'92, 1992, pp. 339–346.

[7] C.M. Newton, Graphics: from alpha to omega in data analysis, in: P.C.C. Wand (Ed.), Graphical Representation of Multivariate Data, Academic Press, New York, 1978, pp. 59–92.

[8] W.S. Cleveland, M.E. McGill, Dynamic Graphics for Statistics, W. Cole, Belmont, CA, 1988.

[9] P. Vellman, The DataDesk Manual, Data Description Inc., Ithaca, NY, 1985.

[10] S.G. Eick, Data visualization sliders, ACM, Proceedings UIST'94, 1994, pp. 119–120.

[11] L. Tweedie, Exploiting interactivity in graphical problem solving: from visual cues to insight, PhD thesis, University of London, 1997.

[12] A. Spoerri, InfoCrystal: a visual tool for information retrieval, IEEE, Proceedings Visualization'93, 1993, pp. 150–157.

[13] L. Tweedie, R. Spence, H. Dawkes, H. Su, Externalizing abstract mathematical models, ACM, Proceedings CHI'96, 1996, pp. 406–412.

Query, Analysis, and Visualization of Hierarchically Structured Data using Polaris

Chris Stolte, Diane Tang, Pat Hanrahan
Stanford University

Abstract

In the last several years, large OLAP databases have become common in a variety of applications such as corporate data warehouses and scientific computing. To support interactive analysis, many of these databases are augmented with hierarchical structures that provide meaningful levels of abstraction that can be leveraged by both the computer and analyst. This hierarchical structure generates many challenges and opportunities in the design of systems for the query, analysis, and visualization of these databases.

In this paper, we present an interactive visual exploration tool that facilitates exploratory analysis of data warehouses with rich hierarchical structure, such as might be stored in data cubes. We base this tool on Polaris, a system for rapidly constructing table-based graphical displays of multidimensional databases. Polaris builds visualizations using an algebraic formalism that is derived from the interface and interpreted as a set of queries to a database. We extend the user interface, algebraic formalism, and generation of data queries in Polaris to expose and take advantage of hierarchical structure. In the resulting system, analysts can navigate through the hierarchical projections of a database, rapidly and incrementally generating visualizations for each projection.

1 Introduction

In the last several years, large OLAP databases have become common in a variety of applications. Corporations are creating large data warehouses of historical data on key aspects of their operations. International research projects such as the Human Genome Project [11] and the Sloan Digital Sky Survey [18] are generating massive scientific databases.

A major challenge with these data warehouses is to extract meaning from the data they contain: to discover structure, find patterns, and derive causal relationships. The sheer size of these data sets complicates this task: Interactive calculations require visiting each record are not plausible, nor is it feasible for an analyst to reason about or view the entire data set at its finest level of detail. Imposing meaningful hierarchical structure on the data warehouse provides levels of abstraction that can be leveraged by both the computer and the analyst.

These hierarchies can come from several different sources. Some hierarchies are known *a priori* and provide semantic meaning for the data. Examples of these hierarchies are Time (day, month, quarter, year) or Location (city, state, country). However, hierarchies can also be automatically derived via data mining algorithms that classify the data, such as decision trees or clustering techniques. Part of the analysis task when dealing with automatically generated hierarchies is in understanding and trusting the results [22].

Visualization is a powerful tool for exploring these large data warehouses, both by itself and coupled with data mining algorithms. However, the task of effectively visualizing large databases imposes significant demands on the human-computer interface to the visualization system. The exploratory process is one of hypothesis, experiment, and discovery. The path of exploration is unpredictable, and analysts need to be able to easily change both the data being displayed and its visual representation. Furthermore, the analyst must be able to first reason about the data at a high level of abstraction, and then rapidly drill down to explore data of interest at a greater level of detail. Thus, the interface must expose the underlying hierarchical structure of the data and support rapid refinement of the visualization.

This paper presents an interactive visual exploration tool that facilitates exploratory analysis of data warehouses with rich hierarchical structure, such as would be stored in data cubes. We base this tool on Polaris [20], a system for the exploration of multidimensional relational databases. Polaris is built upon an algebraic formalism for constructing table-based visualizations. The state of the user interface is a visual specification. This specification is interpreted according to the formalism both to determine the series of queries necessary to retrieve the requested data, as well as determine how to map and layout the resulting tuples into graphical marks. Because every intermediate specification is valid and can be interpreted to create a visualization, analysts can rapidly and incrementally construct complex queries, receiving visual feedback as they assemble and alter the specifications.

The original version of Polaris did not directly support or expose hierarchically structured dimensions, instead presenting each level of the hierarchy as a separate, independent dimension. In this paper, we extend the algebraic formalism (Section 4), user interface (Section 5), and generation of data queries (Section 6) to take advantage of hierarchically structured data cubes. We then illustrate the ease and effectiveness of using Polaris to explore hierarchically structured data via three case studies (Section 7).

2 Related Work

We consider two areas of related work: the visual exploration of databases and the use of data visualization in conjunction with data mining algorithms.

2.1 Visual Exploration of Databases

One area of related work is the field of visual query tools. Projects such as VQE [5], Visage [16], DEVise [14], and Tioga-2 [26] have focused on developing visualization tools that directly support interactive database exploration through visual queries. Users can construct queries and visualizations directly through their interactions with the interface. These systems have flexible mechanisms for mapping query results to graphs and support mapping database tuples to retinal properties of the marks in the graphs. Of these systems, only Tioga-2 provides built-in support for interactively navigating through and exploring data at different levels of detail. However, the underlying hierarchical structure must be created by the analyst during the visualization process; Polaris leverages the hierarchical structure that is already encoded in the data warehouse.

XmdvTool [24], Spotfire [19], and Xgobi [4] provide the analyst with a set of predefined visualizations such as scatterplots and parallel coordinates. These systems are augmented with extensive interaction techniques (e.g., brushing and zooming) that can be used to refine the queries. In contrast, we provide the analyst with a set of building blocks that can be used to interactively construct and refine a wide range of displays to suit the analysis process. Of these systems, only XmdvTool supports the exploration of hierarchically structured data. XmdvTool has been augmented with structure-based brushes [7] that allow the user to control the display's global level of detail (based on a hierarchical clustering of the

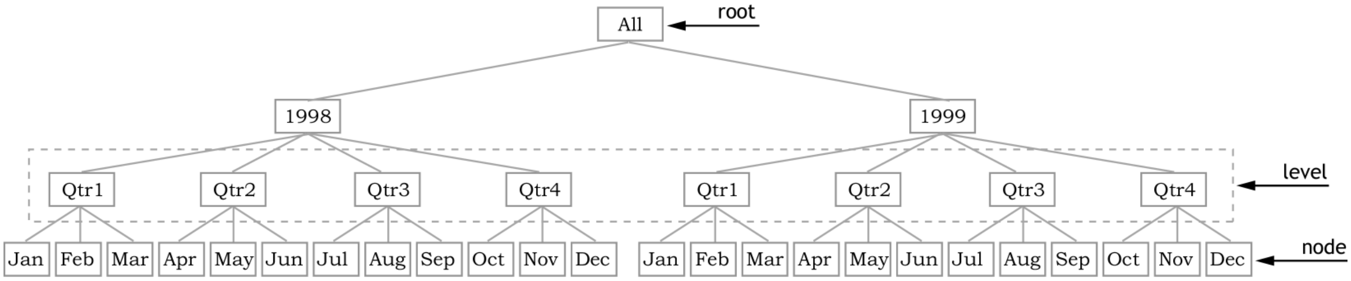

Figure 1: A hierarchical *Time* dimension. A hierarchical dimension is structured as a tree with multiple levels. In this case, there are four levels: *All*, *Year*, *Quarter*, and *Month*. Each level corresponds to a different semantic level of detail. The parent-child relationships in the tree are the basis for aggregation within the dimension.

data) and to brush records based on their proximity within the hierarchical structure. Again, this approach limits the user, in this case to viewing a single hierarchical structuring of the data and a single ordering of that hierarchy to make proximity meaningful. Polaris supports both the simultaneous exploration of multiple hierarchies (derived from semantic meaning or algorithmic analysis) and the ability to reorder the hierarchy as needed.

Another relevant database visualization system is VisDB [12], which focuses on displaying as many data items as possible to provide feedback as users refine their queries. This system even displays tuples that do not satisfy the query, indicating their “distance” from the query criteria using spatial encodings and color. This approach helps the user avoid missing important data points that fall just outside of the selected query parameters. In contrast, Polaris, by taking advantage of the hierarchical structure of the warehouse, provides extensive ability to drill down and roll up data, allowing the analyst to get a complete overview of the data set before focusing on detailed portions of the database.

2.2 Visualization and Data Mining

Many research and commercial systems use visualization in conjunction with automated data mining algorithms. One common application of visualization together with data mining is in helping analysts understand models generated by the data mining process. For example, several researchers have developed techniques specifically for displaying decision trees, Bayesian classifiers, and decision table classifiers [1], and these visualization techniques have been incorporated into products such as SGI’s MineSet [3].

Other approaches to coupling visualization and data mining have traditionally been employed within focused domains. One approach is to use visualization to gain an initial understanding of a warehouse and then apply algorithmic analysis to the identified areas of interest [13][22]. The other major approach is to use data mining to compress the size and dimensionality of the data and then use focused visualization tools to explore the results [10][25].

Unlike these examples, Polaris is not focused on a particular algorithm, a single phase of the discovery process, or a narrow application domain. Instead, Polaris is a general tool that can be used to gain an initial understanding of a warehouse, to visually mine the warehouse, to understand algorithm output, and to interactively explore a mining model. The ability to encode a large number of dimensions in a table layout in Polaris helps an analyst gain an initial understanding of how different dimensions relate as a precursor to automated discovery. Similarly, Polaris can be used directly as a visual mining tool. Finally, by integrating the decision trees and classification networks into the data warehouse as dimensional hierarchies, Polaris can be used by analysts to gain an understanding of how these models classify the data.

3 Background

Polaris [20] was originally designed to support the interactive exploration of multidimensional relational data warehouses rather than data sets with rich hierarchical structure. In this section, we explain the difference between the two types of data sources as well as give a brief overview of Polaris before discussing our extensions to Polaris in the rest of the paper.

3.1 Relational Databases vs. Data Cubes

Relational databases organize data into *relations*, where each row in a relation corresponds to a basic entity or fact and each column represents a property of that entity [23]. For example, a relation may represent transactions in a bank, where each row corresponds to a single transaction, and each transaction has multiple properties, such as the transaction amount transaction, the account balance, the bank branch, and the customer.

We refer to a row in a relation as a *tuple* or *record*, and a column in the relation as a *field*. A single relational database will contain many heterogeneous but interrelated relations.

The fields within a relation can be partitioned into two types: *dimensions* and *measures*. Dimensions and measures are similar to independent and dependent variables in traditional analysis. For example, the bank branch and the customer would be dimensions, while the account balance would be a measure.

In many data warehouses, these multidimensional databases are structured as n-dimensional data cubes. Each dimension in the data cube corresponds to one dimension in the relational schema. Each cell in the data cube contains all the measures in the relational schema corresponding to a unique combination of values for each dimension.

The dimensions within a data cube are often augmented with a hierarchical structure. This hierarchical structure may be derived from the semantic levels of detail within the dimension or generated from classification algorithms. Using these hierarchies, the analyst can explore and analyze the data cube at multiple meaningful levels of aggregation calculated from a base *fact table* (i.e., a relation in the database with the raw data). Each cell in the data cube now corresponds to the measures of the base fact table aggregated to the proper level of detail.

The aggregation levels are determined from the hierarchical dimension, which is structured as a tree with multiple *levels*. Each level corresponds to a different semantic level of detail for that dimension. Within each level of the tree there are many *nodes*, with each node corresponding to a value within the domain of that level of detail of that dimension. The tree forms a set of parent-child relationships between the domain values at each level of detail. These relationships are the basis for aggregation, drill down, and roll up operations within the dimension hierarchy. Figure 1 illustrates the dimension hierarchy for a *Time* dimension.

Simple hierarchies, like the one shown in Figure 1, are commonly modeled using a *star schema*. The entire dimensional hierarchy is represented by a single dimension table (also stored as a relation) joined to the base fact table. In this type of hierarchy, there is only one path of aggregation. However, there are more complex dimension hierarchies where the aggregation path can branch. For

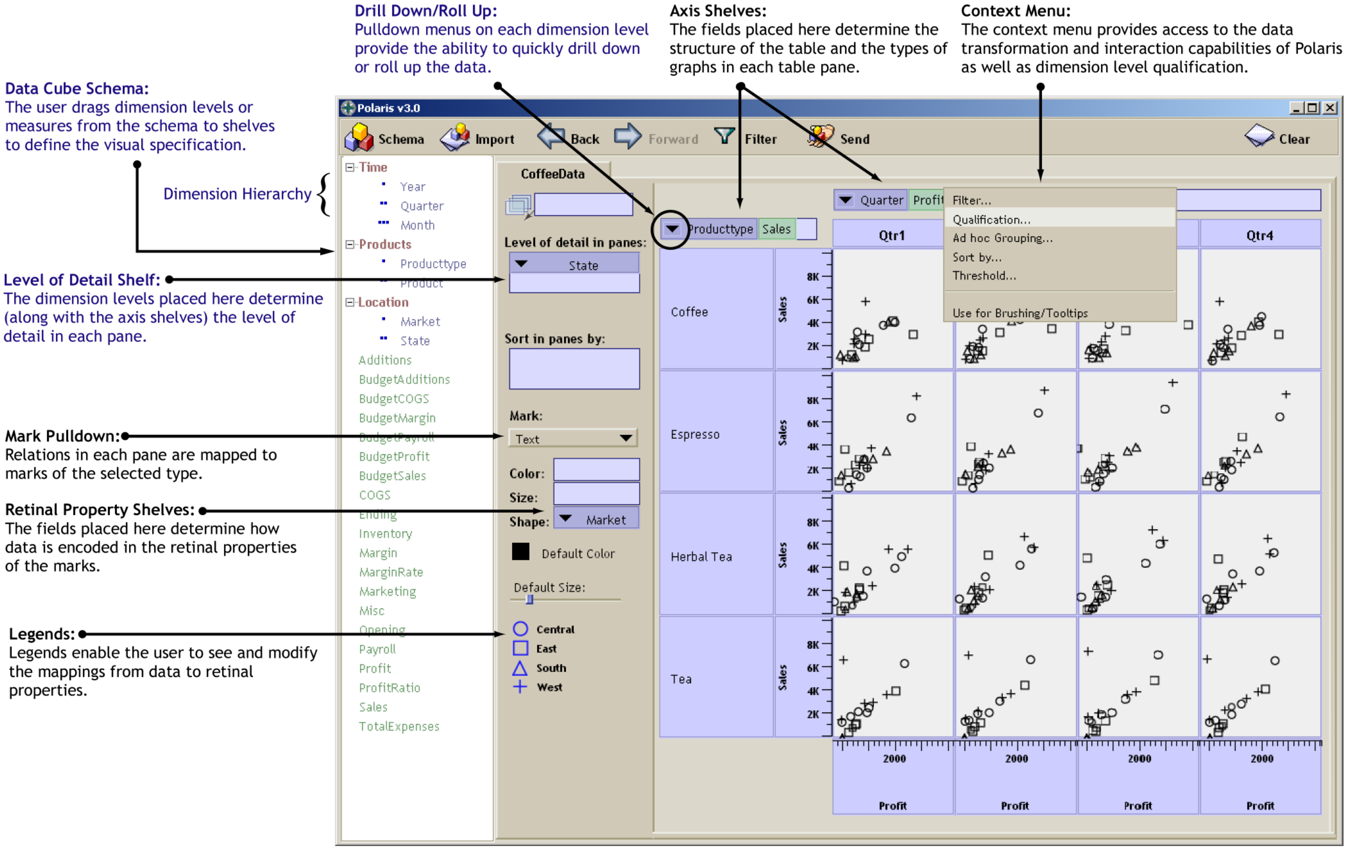

Drill Down/Roll Up:
Pulldown menus on each dimension level provide the ability to quickly drill down or roll up the data.

Axis Shelves:
The fields placed here determine the structure of the table and the types of graphs in each table pane.

Context Menu:
The context menu provides access to the data transformation and interaction capabilities of Polaris as well as dimension level qualification.

Data Cube Schema:
The user drags dimension levels or measures from the schema to shelves to define the visual specification.

Dimension Hierarchy

Level of Detail Shelf:
The dimension levels placed here determine (along with the axis shelves) the level of detail in each pane.

Mark Pulldown:
Relations in each pane are mapped to marks of the selected type.

Retinal Property Shelves:
The fields placed here determine how data is encoded in the retinal properties of the marks.

Legends:
Legends enable the user to see and modify the mappings from data to retinal properties.

Figure 2: The Polaris user interface with enhancements (shown in blue) to expose and support hierarchical dimensions. Analysts construct table-based displays of relational and cube data by dragging dimension levels and measures from the data cube schema onto shelves throughout the display. A given configuration of levels and measures on shelves is called a visual specification. The specification unambiguously defines the analysis and visualization operations to be performed by the system to generate the display.

example, a *Time* dimension might aggregate from *Day* to both *Week* and *Month*. These complex hierarchies are typically represented using a *snowflake schema* that uses multiple relations to represent the diverging hierarchies.

When referring to values within a dimension hierarchy, we will use a dotted notation to specify a specific path from the root level (*All*) of the hierarchy down to the specified value. Specifically, to refer to a value on level m of a hierarchy, we first optionally list the dimension name, then zero or more of the $(m-1)$ intermediate ancestor values, and then finally the value on the m^{th} level, all separated by periods. For example, the *Jan* node on the *Month* level in the *Time* hierarchy that corresponds to January, 1998, can be referred to as *1998.Qtr1.Jan*. When this notation is used, we will call the reference a *qualified value*. When a value is simply described by its node value (without any path to the root node) we call the reference an *unqualified value*.

3.2 Polaris Overview

Before explaining the extensions to Polaris needed for supporting interactive visual exploration of hierarchically structured data sets, we first give a brief overview of the original Polaris system.

The goal of Polaris was to provide an interface for rapidly and incrementally generating table-based displays (note that from here on out, unless otherwise specified as a fact table or dimension table, the term *table* refers to a table-based visualization and not a relation in a database). Users construct these table-based visualizations via a drag-and-drop interface, dragging field names from the Schema box to various blue shelves throughout the interface, as shown in Figure 2. Any configuration of field names on shelves is valid.

The Polaris interface is simple and expressive because it is built upon a formalism for precisely describing graphical table-based visualizations. The configuration of fields on shelves forms a *visual specification*. Each visual specification is an expression of the Polaris formalism that can be interpreted to determine the exact analysis, query, and drawing operations to be performed by the system.

The specification consists of two main portions. The first portion, built on top of an algebra, describes the structure of the table-based visualization (i.e., how the table is divided into panes). We can think of a table as having three axes: the x-axis divides the table into columns, the y-axis divides the table into rows, and the z-axis layers x-y tables that are composited on top of one another. Each intersection of an x, y, and z axis results in a table pane. Thus, the first portion of the specification consists of *table algebra expressions*, with one expression per axis. Each pane contains a set of records (obtained by querying the data cube) that are visually encoded as a set of marks to create a graphic.

While the first portion of the specification determines the "outer table layout," the remaining portion determines the layout within a pane, such as how the data within a pane is transformed for analysis and how it is encoded visually. Specifically, it describes:

1. The sorting and filtering of fields.
2. The mapping of data sources to layers.
3. The grouping of data within a pane and the computation of statistical properties, aggregates, and other derived fields.
4. The type of graphic displayed in each pane of the table. Each graphic consists of a set of marks (e.g., circles, bars, glyphs, etc.), with one mark per record in that pane.
5. The mapping of data fields to retinal properties of the marks in the graphics (for example, mapping *Profit* to the size of a mark or *Quarter* to the color).

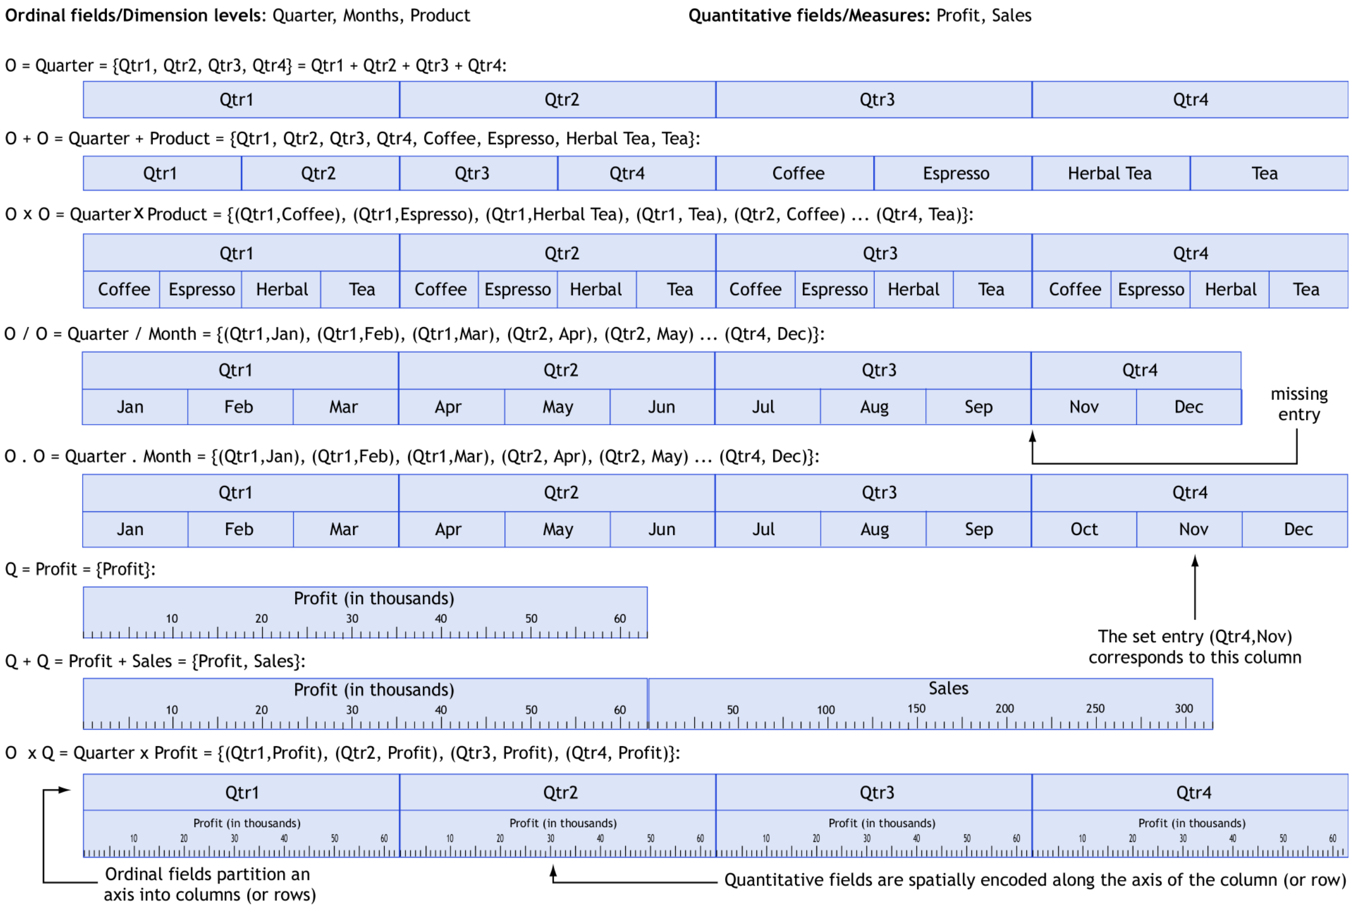
Ordinal fields/Dimension levels: Quarter, Months, Product
Quantitative fields/Measures: Profit, Sales
O = Quarter = {Qtr1, Qtr2, Qtr3, Qtr4} = Qtr1 + Qtr2 + Qtr3 + Qtr4:
Qtr1 | Qtr2 | Qtr3 | Qtr4
O + O = Quarter + Product = {Qtr1, Qtr2, Qtr3, Qtr4, Coffee, Espresso, Herbal Tea, Tea}:
Qtr1 | Qtr2 | Qtr3 | Qtr4 | Coffee | Espresso | Herbal Tea | Tea
O x O = Quarter X Product = {(Qtr1,Coffee), (Qtr1,Espresso), (Qtr1,Herbal Tea), (Qtr1, Tea), (Qtr2, Coffee) ... (Qtr4, Tea)}:
Qtr1 | Qtr2 | Qtr3 | Qtr4
Coffee | Espresso | Herbal | Tea | Coffee | Espresso | Herbal | Tea | Coffee | Espresso | Herbal | Tea | Coffee | Espresso | Herbal | Tea
O / O = Quarter / Month = {(Qtr1,Jan), (Qtr1,Feb), (Qtr1,Mar), (Qtr2, Apr), (Qtr2, May) ... (Qtr4, Dec)}:
Qtr1 | Qtr2 | Qtr3 | Qtr4
Jan | Feb | Mar | Apr | May | Jun | Jul | Aug | Sep | Nov | Dec
missing entry
O . O = Quarter . Month = {(Qtr1,Jan), (Qtr1,Feb), (Qtr1,Mar), (Qtr2, Apr), (Qtr2, May) ... (Qtr4, Dec)}:
Qtr1 | Qtr2 | Qtr3 | Qtr4
Jan | Feb | Mar | Apr | May | Jun | Jul | Aug | Sep | Oct | Nov | Dec
Q = Profit = {Profit}:
Profit (in thousands)
10 20 30 40 50 60
The set entry (Qtr4,Nov) corresponds to this column
Q + Q = Profit + Sales = {Profit, Sales}:
Profit (in thousands)
10 20 30 40 50 60
Sales
50 100 150 200 250 300
O x Q = Quarter x Profit = {(Qtr1,Profit), (Qtr2, Profit), (Qtr3, Profit), (Qtr4, Profit)}:
Qtr1 | Qtr2 | Qtr3 | Qtr4
Profit (in thousands)
10 20 30 40 50 60
Ordinal fields partition an axis into columns (or rows)
Quantitative fields are spatially encoded along the axis of the column (or row)

Figure 3: Example of the set interpretations and table structures resulting from simple applications of the table algebra operators. Ordinal fields (e.g., dimension levels) partition the table into columns (or rows) and quantitative fields (e.g., measures) are spatially encoded as axes within the columns. Note the difference between the application of the nest and dot operator to the same operands when the fact table does not contain data for October.

We only need to extend the table algebra and the specification of the filtering and sorting. The rest of the formalism, including how we determine the type of graphic and the visual encodings, has not changed and so we do not discuss them further here. See Stolte et al. [20] for a detailed discussion.

Thus, in order to extend Polaris to support hierarchical dimensions, we need to modify:

- the formalism and table algebra (described in Section 4)
- the user interface (described in Section 5)
- and the interpretation of the visual specification as a set of queries in a multidimensional query language (described in Section 6).

4 Extending the Formalism

In order to support both relational databases and hierarchically structured data cubes, we need to extend two aspects of the Polaris formalism: the specification of the table configurations, and the filtering and sorting of fields. Before we discuss these two extensions, however, we first give a brief review of the table algebra.

4.1 Table Algebra Review

A key component of this formalism is the table algebra, which is used to specify the table configurations. When analysts place fields on the axis shelves (shown in Figure 2) they are implicitly creating expressions in this algebra. A complete table configuration consists of three separate expressions. Two of the expressions define the configuration of the x and y axes of the table, partitioning the table into rows and columns. The third expression defines the z axis of the table, which partitions the display into layers.

Each expression is composed of operands connected by operators. Each operand is evaluated to a set form, and the operators define how to combine two sets. Thus, each expression can be interpreted as a single set (the *normalized set form*), where each element in the set corresponds to a single row, column, or layer.

To be more specific, each operand of the table algebra is the name of a field. There are two types of operands: ordinal and quantitative. Whether an operand is ordinal or quantitative depends on the type of the corresponding field in the database.

The set interpretation of an ordinal operand consists of the members of the ordered domain of the field. For example, the set interpretation of the *Month* operand would be $\{Jan, Feb, \ldots, Dec\}$. The set interpretation of a quantitative operand is a single-element set containing the field name. For example, the set interpretation of the Profit operand would be $\{Profit\}$.

The assignment of sets to the different types of operands reflects the difference in how the two types of fields are encoded into the structure of the table. Ordinal fields partition the table into rows and columns, whereas quantitative fields are spatially encoded as axes within the table panes. Examples of the set interpretations and resulting table structures for both ordinal and quantitative operands are shown in Figure 3.

As stated above, a valid expression in the algebra is an ordered sequence of one or more operands with operators between each pair of adjacent operands. The operators in this algebra, in order of precedence, are cross (×), nest (/), and concatenation (+); paren-

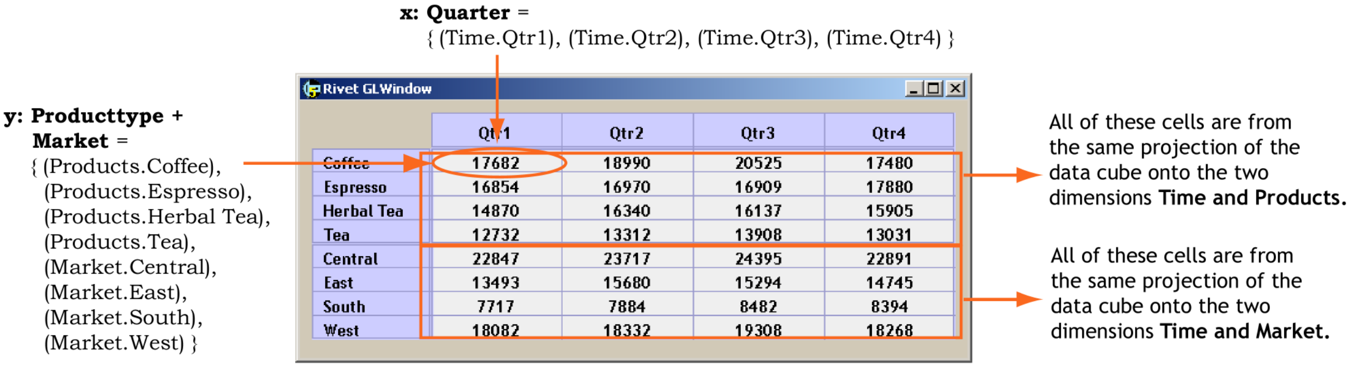

x: Quarter =
{ (Time.Qtr1), (Time.Qtr2), (Time.Qtr3), (Time.Qtr4) }

y: Producttype +
Market =
{ (Products.Coffee),
(Products.Espresso),
(Products.Herbal Tea),
(Products.Tea),
(Market.Central),
(Market.East),
(Market.South),
(Market.West) }

Rivet GLWindow

	Qtr1	Qtr2	Qtr3	Qtr4
Coffee	17682	18990	20525	17480
Espresso	16854	16970	16909	17880
Herbal Tea	14870	16340	16137	15905
Tea	12732	13312	13908	13031
Central	22847	23717	24395	22891
East	13493	15680	15294	14745
South	7717	7884	8482	8394
West	18082	18332	19308	18268

All of these cells are from the same projection of the data cube onto the two dimensions **Time and Products.**

All of these cells are from the same projection of the data cube onto the two dimensions **Time and Market.**

Figure 4: Each pane in a Polaris visualization corresponds to a slice of a projection of a data cube. The projection in each pane is determined by the contents of the "Level of Detail" shelf and by the normalized set form of the table expressions. The table is partitioned into rows, columns, and layers corresponding to the entries in these sets. The underlying data cube must be projected to include only the dimensions that occur in these entries. This is shown here for a simple text-based table. Generating this table requires two separate projections of the data cube because of the concatentation in the y-axis expression.

theses can be used to alter the precedence. Because each operand is interpreted as an ordered set, the precise semantics of each operator are defined in terms of how they combine two sets (one each from the left and right operands) into a single set. Some examples are shown in Figure 3.

Thus, every expression in the algebra can be reduced to a single set, with each entry in the set being an ordered concatenation of zero or more ordinal values followed by zero or more quantitative field names. For example, the normalized set form of the expression *Month* × *Profit* is { (*Jan*, *Profit*), (*Feb*, *Profit*), . . . , (*Dec*, *Profit*) }. The normalized set form of an expression determines one axis of the table: the table axis is partitioned into columns (or rows or layers) so that there is a one-to-one correspondence between columns and entries in the normalized set.

4.2 Redefining the Algebra Operands

In order to fully support and expose the hierarchical structure in data cube dimensions, we must redefine the algebra so that the operands are measures and dimension levels rather than independent database fields. In this redefinition, measure operands are trivially treated the same as quantitative fields: we assign to each measure operand a single element set containing the measure name. Like quantitative fields in the original algebra, measures will be spatially encoded as axes within the panes.

Similarly, we would like to treat dimension levels in the same way we treated ordinal fields and assign to each the ordered domain of the dimension level. The resulting sets would then partition the table into rows, columns, and layers. There are, however, complications. The domain of a dimension level is not a single ordered list. Instead, it is composed of the node values at a particular level in the dimension hierarchy, and each node value is uniquely defined by its path to the root of the hierarchy. To illustrate the complications this causes, we consider the *Time* hierarchy illustrated in Figure 1.

First, consider the *Month* level of the hierarchy. One possible set interpretation of this symbol would be to list each node value, including its path to the root for uniqueness, ordered by a depth-first traversal of the dimension hierarchy; e.g., {*1998.Qtr1.Jan*, . . . , *1999.Qtr4.Dec*}. Although this approach provides a unique set interpretation for each dimension level, it limits the expressiveness of the algebra. Any table constructed to include *Month* must also include *Year*; it is not possible to create displays that summarize monthly values across years, a useful view that we would like to support. Interestingly, however, summarizing monthly values across years is not a standard projection of a data cube, as it requires aggregating across a hierarchical level. We discuss how this type of aggregation is computed in Section 6.

A second approach would be to list only the node values, ignoring the path to the root of the hierarchy and excluding repeated values. Again, we order the node values by a depth-first traversal of the dimension hierarchy. For *Month*, this would yield {*Jan*, *Feb*, . . . , *Nov*, *Dec*}. Clearly, using this set interpretation we can generate displays that summarize monthly values across years. Furthermore, we can generate displays that drill down into a hierarchy by using our nest ("/") operator; e.g., *Year* / *Month*.

The use of the nest for drilling down into a hierarchy, however, would be flawed. The nest operator is unaware of the defined hierarchical relationship between the dimension levels but instead works by deriving a relationship based on the tuples in the fact table. Not only is this inefficient, as fact tables are often quite large, but it can also yield incorrect results. For example, consider the situation where no data was logged for October. Application of the nest operator would result in an incorrectly derived *Time* hierarchy that did not include October as a child of Qtr4 or either year (see Figure 3).

Our solution is to introduce another operator, the dot (".") operator, that is similar to the nest operator but "hierarchy-aware." We review the definition of nest and then define dot. If we define *FT* to be the fact table being analyzed, *r* to be a record, and *A(r)* to be the value of the field *A* for the record *r*, then the definition of nest, as presented in [20], is:

$$A/B = \{(a, b) \mid \exists r \in FT \ st \ A(r) = a \ \& \ B(r) = b\}$$

The dot operator is defined similarly. If we define *DT* to be the relational dimension table defining the hierarchy that contains the levels *A* and *B*, and *A* precedes *B* in the schema of *DT*, then:

$$A.B = \{(a.b) \mid \exists r \in DT \ st \ A(r) = a \ \& \ B(r) = b\}$$

Note that whereas nest produces a set of two-valued tuples, dot produces a set of single-valued tuples, each containing a qualified value. If the two operands are not levels of the same dimension hierarchy (or set interpretations of operations on levels of the same hierarchy), or *A* does not precede *B* in the schema of *DT* (e.g., *A* must be an ancestor level in the tree defined by *DT*) , then the dot operator evaluates to the empty set. With this definition, the two expressions *Month* and *Year.Month* are not equivalent: *Month* is interpreted as {*Jan*, *Feb*, . . . , *Dec*} whereas *Year.Month* would be interpreted as {*1998.Jan*, *1998.Feb*, . . . , *1999.Dec*}. With a fully populated fact table, *Year.Month* is equivalent to *Year* / *Month*.

Given these set interpretations for dimension and measure operands, we can apply the set semantics for each operator to reduce expressions in this new algebra to their normalized set form, with each entry in the normalized set being an ordered concatenation of zero or more domain values followed by zero or more measure names. As before, the normalized set form determines one axis of the table.

4.3 Filtering and Sorting within the Algebra

In our original formalism, a table configuration was specified by three expressions in the table algebra, and then filtering and sorting was specified separately by listing the sorted and filtered domain for each database field that was to be filtered or sorted. When the set interpretation was generated for field operands in the algebra, these specified domains would be used. It is possible, however, to generate a more succinct and general formalism if we incorporate the filtering and sorting directly into the table algebra.

In our revised formalism, if a dimension or measure is to be filtered (or sorted), then the filtered and sorted domain is listed directly after the instance of the level or measure operand in the expression, in effect directly specifying a set interpretation for the operand. For example, if we wished to filter the expression *Month* + *ProductType* to include only the first three months of the year, sorted in reverse order, this could be specified by including the filtered domain in the expression as follows: *Month*{*Mar*, *Feb*, *Jan*} + *ProductType*. The advantage provided by this revision of the table algebra is the ability to specify separate filters and orderings for different instances of the same operand in an expression. Similarly, we can filter a measure by specifying a range of values, e.g., *Profit*{0, 500}.

We also need to allow the use of qualified values in the specification of filtering or sorting of dimension levels. As we discussed in Section 3.1, a value in a dimension hierarchy can either be described by simply stating the value in the node (an unqualified value) or by describing a path from that node to the root node in the hierarchy (a qualified value). When filtering or sorting a dimension level, it is necessary to be able to use both types of values in the specification, as the unqualified node values are often not unique. For example, if the user wishes to exclude *1998.Jan* but not *1999.Jan*, then qualified values must be used.

5 Redefining the User Interface

Having redefined the formalism underlying the Polaris interface, we must now alter the interface to support hierarchically structured data. Five major changes need to be made:

1. the Schema list must display dimension hierarchies and measures, not simply database fields;
2. the analyst must be able to distinguish between *Month* and *Year.Month* when including *Month* in a specification;
3. the analyst must be able to filter a dimension level using qualified values;
4. the analyst must be able to quickly drill down and roll up a dimension hierarchy using the interface;
5. the analyst needs to be able to change the number of marks within each pane to reflect different levels of detail.

Figure 2 illustrates the revised interface. We now discuss each interface extension in detail.

5.1 The Schema

In the original interface, the analyst was presented with a list box containing the ordinal and quantitative fields in the database. The analysts included these fields in a specification simply by dragging and dropping the field's name onto the appropriate shelf in the interface. To support hierarchical data cubes, we have extended this list box to display the dimensions of the data cube with an ordered list of the dimension's levels beneath each dimension. The analysts can drag and drop any dimension level to the interface as they did with the ordinal fields of the database. The dimension's name, however, cannot be dragged to the interface; the analyst can only manipulate the individual levels within a dimension.

5.2 Qualifying Dimension Levels

When an analyst drops a dimension level, such as *Month*, on a shelf, there are several potential intentions. He may intend to include the operand *Month* in an expression, but he may also mean *Year.Month* or *Year.Quarter.Month*; the analyst needs to be able to specify the exact qualification desired. Our solution is the make full qualification (e.g., *Year.Quarter.Month*) the default. To generate a different qualification, the user can right-click the dimension level in the shelf and select the "Qualification..." menu item. He is then presented with a dialog box that allows him to explicitly specify which of the intermediate levels to include in the qualification of the operand, thus generating the applicable expression.

5.3 Qualifying Dimension Level Filters

When applying filters to a dimension level, an analyst may want to specify the filter using either qualified or unqualified values. We have extended the Polaris interface to allow both options. For example, if the user wishes to exclude *1998.Jan* but not *1999.Jan*, he can choose to filter using qualified values. Similarly, it is possible to specify a filtering using unqualified values: each qualified value that matches the unqualified value will be included in the filter. Currently, Polaris requires the filter be specified using either qualified or unqualified values, but not both. As a future extension, we intend to support heterogeneous filtering.

5.4 Drilling down and Rolling up

When analyzing and exploring large data cubes, a common operation is to drill down or roll up within a dimension hierarchy. Therefore, it is important to include a simple mechanism for performing these operations. One option is for the analyst to remove the current level from the appropriate shelf (by dragging it off the shelf) and then drag the new level to that same shelf. Although the desired effect is achieved, it is more complicated than we would like.

We provide an alternate mechanism for drilling down and rolling up a dimension. Within the box representing each dimension level on a shelf, there is an "∇" icon, as can be seen in Figure 2. When the user clicks on the "∇" icon, he is presented with a listing of all the levels of the dimension (including diverging levels in complex dimensional hierarchies). When a new level is selected, this is interpreted as a drill down (or roll up) operation along that dimension and the current level is automatically replaced with the selected level (with the same qualification). Thus, the user can rapidly move between different levels of detail along a dimension, refining the visualization as he navigates.

5.5 Grouping within Panes

In the original version of Polaris, the analyst specified the grouping of tuples within each pane by placing fields on the shelf titled "Group By". Each field in this shelf was included in the GROUP BY clause in the SQL query that aggregated the data in each pane into tuples to be mapped to marks.

The situation when visualizing data cubes is slightly different. The query for each pane does not produce a relational data set that is then grouped and aggregated. Instead, each pane corresponds to a projection of the data cube, with the projection determined by the dimension levels included in the table expressions. To produce additional marks within a pane, the analyst must specify additional dimensions to be included in these projections, done by including the desired dimension levels in the "Level of Detail" shelf (shown in Figure 2), the hierarchical analog of the "Group by" shelf.

As was the case with the original version of Polaris, this "Level of Detail" shelf gives the analyst the ability to rapidly drill down

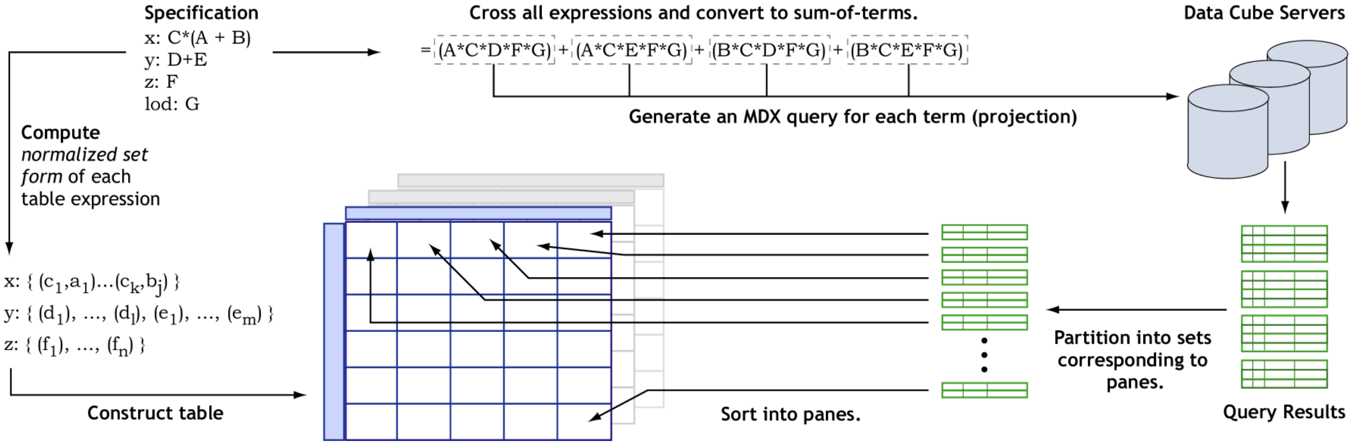
Specification
x: C*(A + B)
y: D+E
z: F
lod: G
Cross all expressions and convert to sum-of-terms.
= (A*C*D*F*G) + (A*C*E*F*G) + (B*C*D*F*G) + (B*C*E*F*G)
Data Cube Servers
Generate an MDX query for each term (projection)
Compute *normalized set form* of each table expression
x: { (c1,a1)...(ck,bj) }
y: { (d1), ..., (dl), (e1), ..., (em) }
z: { (f1), ..., (fn) }
Construct table
Sort into panes.
Partition into sets corresponding to panes.
Query Results

Figure 5: The overall data flow in Polaris when generating a visual representation of a data cube.

into their data without changing the table configuration. Changing the level of detail without changing the table configuration only changes the data density within each pane.

6 Querying a Hierarchical Data Cube

The final step in extending Polaris to fully support hierarchical data cubes is to show how to construct an efficient set of multidimensional queries from a specification in our formalism.

Each pane in a Polaris visualization corresponds to either a slice of a projection of the data cube or an aggregation of such a projection. The specific projection corresponding to each pane is determined by the contents of the “Level of Detail” shelf (discussed in Section 5.5) and by the normalized set form of the table axis expressions (discussed in Section 4). The table is partitioned into rows, columns, and layers corresponding to the entries in these sets. Therefore, each pane in the table is associated with three set entries corresponding to its row, column, and layer, respectively.

The underlying data cube must be projected to include only the dimension levels that occur in the three set entries (and in the Level of Detail shelf) and it must be sliced to include only the specific dimension members that occur in these entries. This is illustrated in Figure 4. When multiple set entries defining a pane refer to different levels of the same dimension, the correct projection to retrieve is the one corresponding to the most detailed level of that dimension. Before determining how the projections are efficiently retrieved from the server, we must carefully consider the situation where set entries contain values whose qualification skips levels in the hierarchy.

When set entries contain values whose qualification skips levels (e.g., *Time.Jan*), this is interpreted to imply that nodes in the hierarchy whose values are not unique (when we consider only the included levels) should be aggregated in that projection. In the *Time.Jan* example, the aggregation for the pane must be computed by aggregating across years, and thus across a hierarchy level rather than up the hierarchy. This type of aggregation is not natively supported in most hierarchical query languages. Thus, we request the cube projection from the remote server and compute the aggregation within Polaris before sorting tuples into panes, as shown in Figure 4. If all node values are unique across the entire level, then no aggregation needs to be performed.

Although it is possible for each pane to correspond to a different projection of the cube, the common situation is for a large number of panes to correspond to the same projection and differ only by how that projection is sliced. For efficiency, we would like to consider these panes as a group and send a single query to the OLAP server requesting the appropriate projection (and then, if necessary, perform a single aggregation of the projection). The projection can then be sorted into panes locally.

The key to efficiently utilizing the OLAP server is this grouping of queries. By algebraically manipulating our table expressions, we can quickly determine all projections corresponding to a given table configuration. The key observation is that of our four algebraic operators (nest, cross, concatenate, and dot), the only operator that can produce adjacent panes with differing projections is the concatenate operator. Nest, cross, and dot include all input dimension levels in each output set entry; concatenate does not. Thus, if we compute a single expression as the cross of the three table expressions and then reduce to a sum-of-terms form, the resulting terms will correspond to the set of projections that need to be generated. This is illustrated in Figure 5.

Most typical multidimensional query languages provide a mechanism for generating projections of the data cube. Our current implementation generates a single MDX query to a remote Microsoft Analysis Server for each projection. The resulting cells are then sorted into panes using transformation capabilities built into Polaris. In addition, any explicitly specified filtering of dimension members is included in the MDX queries sent to the remote server. The overall data flow in Polaris is depicted in Figure 5.

7 Results

In this section, we illustrate how Polaris can be used to effectively navigate and analyze three hierarchically structured data sets: (1) a 12-week trace of mobile network usage, (2) results from the 2000 presidential election, and (3) historical business metrics for a hypothetical coffee chain.

7.1 Mobile Network Usage Data

Figure 6 shows an analysis of a 12-week trace of every packet that entered or exited the mobile network in the Gates building at Stanford University [21]. Over the 12 weeks, 78 million packet headers were collected. The analysis goal is to understand usage patterns of the mobile network. This data is stored in a data cube with many different dimensions (User, Time, Remote host, Traffic direction, and Application), each with multiple levels of detail.

To start the analysis, the analyst first sees if she can spot any patterns in time, so she creates a series of line charts in Figure 6(a) showing packet count and size versus time for the most common applications, broken down and colored by the direction of the traffic. In these charts, the analyst can see that the web is the most consistently used application, while session is almost as consistent. File transfer is the least consistent, but also has some of the highest peaks in both incoming and outgoing ftp traffic. Note the log scale on the y-axes.

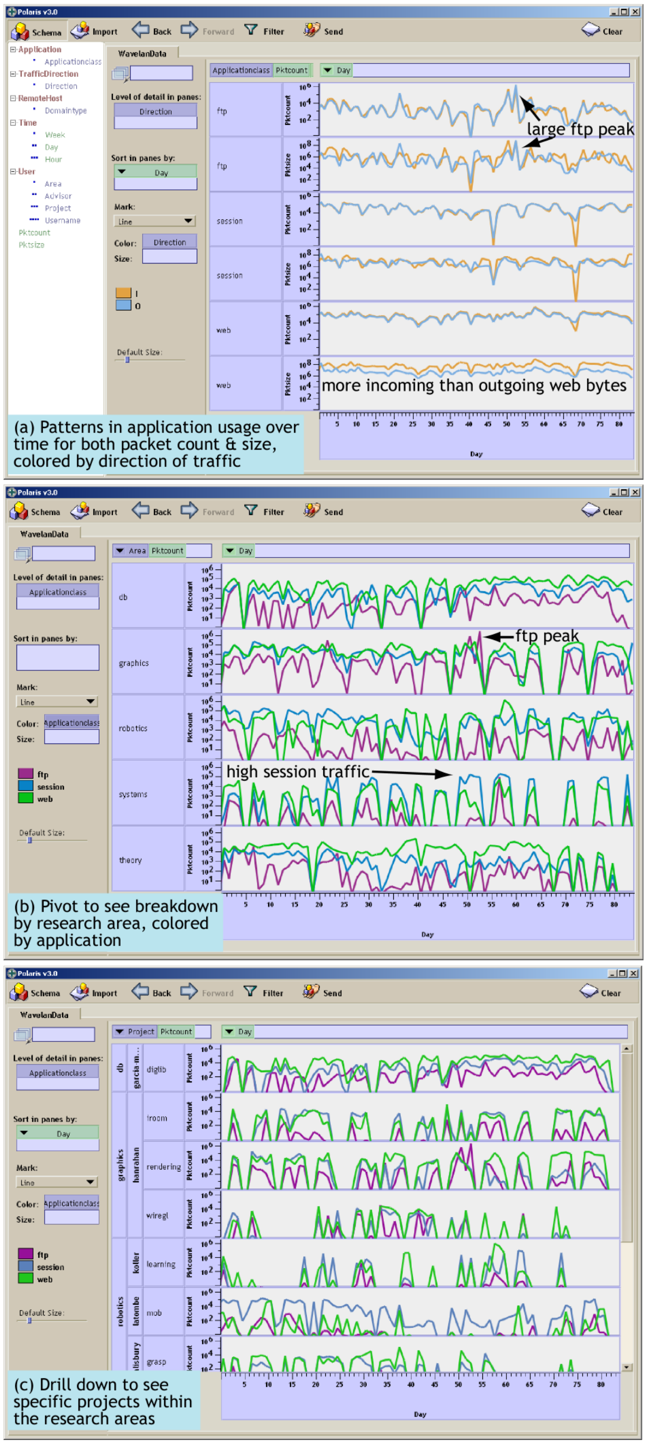
(a) Patterns in application usage over time for both packet count & size, colored by direction of traffic

large ftp peak

more incoming than outgoing web bytes

(b) Pivot to see breakdown by research area, colored by application

ftp peak

high session traffic

(c) Drill down to see specific projects within the research areas

Figure 6: Analysis of network usage data using Polaris.

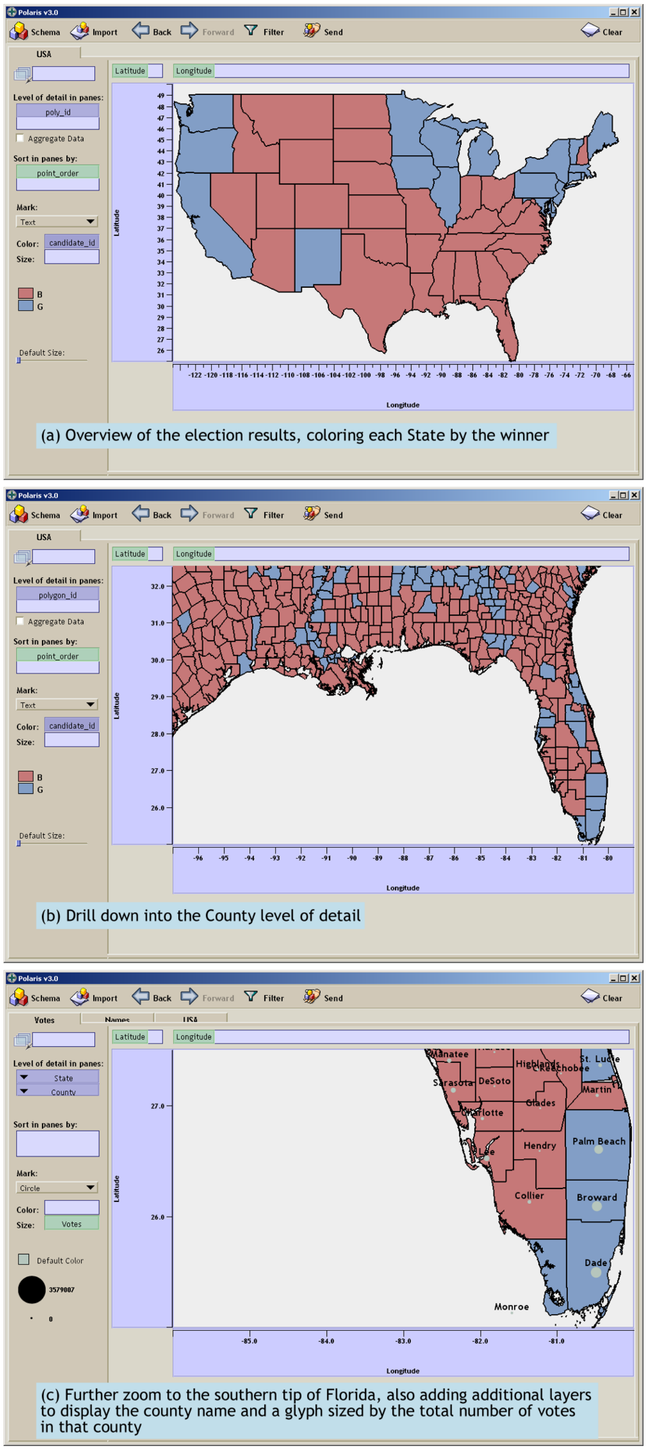
(a) Overview of the election results, coloring each State by the winner

(b) Drill down into the County level of detail

(c) Further zoom to the southern tip of Florida, also adding additional layers to display the county name and a glyph sized by the total number of votes in that county

Figure 7: Analysis of the results of the 2000 presidential election.

Given this broad understanding of traffic patterns, the next question posed by the analyst is how the application mix varies depending on the research area. The analyst pivots the display to generate a single line chart of packet count per research area over time, broken down and colored by application class (Figure 6(b)). From this breakdown, the analyst can see that the graphics group was responsible for the large incoming and outgoing file transfers. She can also see that the systems group had atypically high session traffic.

Curious, the analyst then drills down further to see the individual project groups (Figure 6(c)), discovering that the large file transfers were due to the rendering group within the graphics lab, while the robotics lab had vastly different behavior depending on the particular group (the mob group dominated by session traffic, while the learning group had more web traffic, for example).

7.2 2000 Presidential Election Results

Figure 7 shows Polaris being used to explore and analyze the results of the 2000 presidential election. This data is particularly interesting because the visualizations used to explore it are created from two separate data sets. The first data set is a relational database

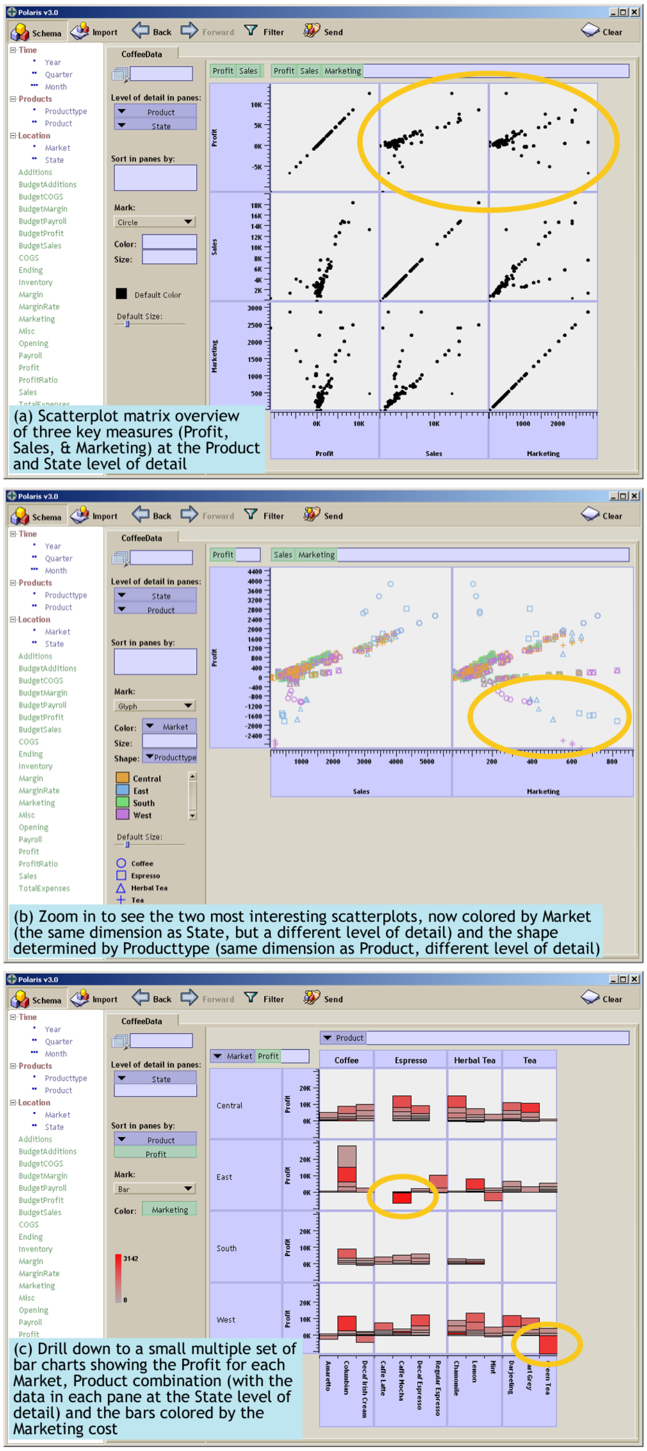

(a) Scatterplot matrix overview of three key measures (Profit, Sales, & Marketing) at the Product and State level of detail

(b) Zoom in to see the two most interesting scatterplots, now colored by Market (the same dimension as State, but a different level of detail) and the shape determined by Producttype (same dimension as Product, different level of detail)

(c) Drill down to a small multiple set of bar charts showing the Profit for each Market, Product combination (with the data in each pane at the State level of detail) and the bars colored by the Marketing cost

Figure 8: Analysis of sales data for a hypothetical coffee chain.

of approximately 500,000 tuples (stored in Microsoft's SQLServer) describing detailed polygonal outlines of the states and counties in the USA. Additional levels of detail have been constructed by polygon simplification, and the resulting levels of detail form a Location hierarchy. The second data set is stored as a data cube (in Microsoft's Analysis Server) and contains detailed county-by-county vote results (also with a Location dimension). In the first two visualizations, these data sets are explicitly joined before being imported into Polaris. In the final visualization, we use the ability in Polaris to visually join data sets using layers.

In Figure 7(a), the analyst has generated an overview of the entire country at the State level in the Location hierarchy, coloring each state by which candidate won that state. The analyst is interested in more detailed results for the state of Florida, so she filters on the Latitude and Longitude measures to focus on Florida and changes the level of detail to County, generating Figure 7(b). In the final visualization, shown in Figure 7(c), the analyst further focuses on the southern tip of Florida (by again filtering the Latitude and Longitude measures). Furthermore, she adds two additional layers to the visualization (read directly from the data cube) and displays both the name and the total number of votes counted in each county.

7.3 Historical Profit/Sales for a Coffee Chain

The final analysis is shown in Figure 8. The data being analyzed is two years of business metrics for a hypothetical nationwide coffee chain, comprising approximately 5,000 tuples stored in a data cube. The data is characterized by three main dimensions (Time, Products, and Location), each with multiple levels of detail. We consider a scenario where the analyst is concerned with reducing marketing expenses and is trying to identify products that are not generating profit and sales proportional to their marketing costs.

The first visualization created, Figure 8(a), is an overview of three key measures (Profit, Sales, and Marketing) as a scatterplot matrix. The analyst has drilled down using the Level of Detail shelf to the Product and State level. The two charts circled in orange show that several of the distributions do not reflect the positive correlations that the analyst was expecting. To further investigate, the analyst reduces the scatterplot matrix to two graphs and colors the records by Market and Producttype (Figure 8(b)), thus identifying espresso products in the East region and tea products in the West region as having the worst marketing cost to profit ratios.

In the final visualization, Figure 8(c), the analyst drills down into the data to get a more detailed understanding of the correlations: She creates a small multiple set of stacked bar charts, one for each Market and Producttype. Within each chart, the data is further drilled down by individual Product and State. Finally, each bar is colored by the Marketing cost. As can be seen in the visualization, several products such as Caffe Mocha in the East have negative profit (a descending bar) with high marketing cost (a bright red bar). Having identified such poorly performing products, the analyst can modify the marketing costs allocated to them.

7.4 Summary

Each of these case studies demonstrates how analysis progresses from a high level of abstraction to detailed views of the data. Furthermore, each example shows the importance of being able to easily change the data being viewed, pivot dimensions, and drill down during the analysis process.

8 Discussion

In this section, we focus on two points of discussion. First, we discuss the different roles Polaris can play in the knowledge discovery process, and second, we discuss how our formalism can be applied to the development of generalized visualization systems, particularly level of detail systems.

In this paper, we have demonstrated the effectiveness of Polaris as a stand-alone tool for visual mining of large, hierarchical databases. Equally important is how Polaris can be coupled with automated data mining systems to help analysts better understand not only their data, but also the models generated by the algorithms. First, Polaris can be used as a precursor to data mining: The analyst benefits from an understanding of the overall structure of the data that helps her steer the discovery process and provides context for "hidden information" discovered by the algorithms. Second, Polaris can also be used to validate and comprehend the models and results generated by algorithmic analysis. Analysts do not want to treat an algorithm as a black box and blindly trust its output. One technique for using Polaris for validation is to construct hierarchical dimensions from the output generated by classification algorithms. The analyst can then drill down and roll up the data, traversing the classification hierarchy and inspecting the records sorted into each bucket, further developing understanding and trust.

A second point of discussion is the application of our formalism to the development of general visualization systems. Although we have only demonstrated our formal language as an underlying technology for the Polaris interface, we believe it is a promising basis for the development of a wide-range of visualization systems. One example is in the development of interactive "semantic-zooming" visualization systems. Programmers developing such systems need a mechanism for describing a wide range of visual displays, with each display being associated with a different level of detail view of the data. Using our formalism, these programmers could simply describe each visual display with a succinct specification. When the user interacts with the interface to move to a different level of detail, the system need only feed the appropriate specification into our interpreter. The interpreter would generate all of the drawing operations and queries necessary to generate the display. In addition to simplifying the development of such systems, the presence of an underlying formalism also serves to help clearly define the semantics of the interface, as demonstrated by Polaris.

9 Conclusions and Future Work

We have extended Polaris, an interface for the exploration and analysis of large multidimensional databases, to fully support and expose the hierarchical structure of data cubes. These dimensional hierarchies play an indispensable role in the analysis of large databases where, in order for the analysis task to be manageable, it is necessary to perform the analysis at multiple levels of aggregation, moving from visual overviews to details on demand. In extending Polaris, we have extended not only the interface, but also the underlying algebraic formalism. Furthermore, we have developed an efficient mechanism for interpreting the formal specifications as a collection of multidimensional queries.

We have many plans for future work in extending this system. As databases continue to grow in size, developing tools and techniques for the interactive exploration of data at multiple levels of detail is crucial. As we discussed in Section 8, we believe our algebraic formalism provides a solid foundation upon which to build visualization systems. We are currently building systems that, a la Pad++ [17], automatically and interactively change the visual representation as the analyst changes level of detail. This research has many interesting challenges, including transitions between different visual representations, mapping representations to levels of detail, and maintaining interactivity while exploring large data warehouses.

A second area of future research is the visual presentation of metadata. Hierarchically structured dimensions are one instance of an increasingly popular trend: the augmentation of data with rich domain-specific metadata. This metadata is as important to the analysis process as the underlying database itself. An important area of future research is the development of visualization techniques that display this metadata and leverage it in the display of the described data.

References

[1] B. Becker. Visualizing Decision Table Classifiers. In Proc. of Information Visualization, October, 1998, pp. 102-105.

[2] B. Becker, R. Kohavi, and D. Sommerfield. Visualizing the Simple Bayesian Classifier. In KDD Workshop on Issues in the Integration of Data Mining and Data Visualization, 1997.

[3] C. Brunk, J. Kelly, and R. Kohavi. MineSet: an integrated system for data mining. In Proc. of the 3rd International Conference on Knowledge Discovery and Data Mining, 1997, pp. 135-138.

[4] A. Buja, D. Cook, and D. F. Swayne. Interactive High-Dimensional Data Visualization. In Journal of Computational and Graphical Statistics, 5(1), 1996, pp. 78-99.

[5] M. Derthick, J. Kolojejchick and S. F. Roth. An Interactive Visualization Environment for Data Exploration. In Proc. of Knowledge Discovery in Databases, August, 1997, pp. 2-9.

[6] S. Eick. Visualizing Multi-Dimensional Data. In Computer Graphics, February 2000, pp. 61-67.

[7] Y. Fua, M. O. Ward, and E. Rundensteiner. Navigating Hierarchies with Structure-Based Brushes. In Proc. of Information Visualization, October 1999, pp. 58-64.

[8] J. Goldstein, S. F. Roth, J. Kolojejchick, and J. Mattis. A Framework for Knowledge-based Interactive Data Exploration. In Journal of Visual Languages and Computing, December 1994, pp. 339-363.

[9] J. Gray, S. Chaudhuri, A. Bosworth, A. Layman, D. Reichart, M. Venkatrao, H. Pirahesh, and F. Pellow. Data Cube: A Relational Aggregation Operator Generalizing Group-By, Cross-Tab, and Sub-Totals. In Proc. of the 12th International Conference on Data Engineering, February 1996, pp. 152-159.

[10] C. G. Healey. On the Use of Perceptual Cues and Data Mining For Effective Visualization of Scientific Datasets. In Proc. Graphics Interface, 1998, pp. 177-184.

[11] Human Genome Project. [online] Available: http://www.ornl.gov/hgmis/about.html, cited February 2002.

[12] D. Keim and H.P. Kriegel. VisDB: Database Exploration using Multidimensional Visualization. In IEEE Computer Graphics and Applications, 14(5), 1994, pp. 40-49.

[13] R. Kohavi. Data Mining and Visualization. Invited talk at the National Academy of Engineering US Frontiers of Engineers, Sept 2000. Available in Frontiers of Engineering: Reports on Leading-Edge Engineering From the 2000 NAE Symposium on Frontiers of Engineering, National Academy Press, 2001.

[14] M. Livny, R. Ramakrishnan, K. Beyer, G. Chen, D. Donjerkovic, S. Lawande, J. Myllymaki and K. Wenger. DEVise: Integrated Querying and Visual Exploration of Large Datasets. In Proc. of ACM SIGMOD, May, 1997.

[15] S.F. Roth, J. Kolojejchick, J. Mattis and J. Goldstein. Interactive Graphic Design Using Automatic Presentation Knowledge. In Proc. of SIGCHI '94, April 1994, pp. 112-117.

[16] S.F. Roth, P. Lucas, J.A. Senn, C.C. Gomberg, M.B. Burks, P.J. Stroffolino, J. Kolojejchick and C. Dunmire. Visage: A User Interface Environment for Exploring Information. In Proc. of Information Visualization, October 1996, pp. 3-12.

[17] K. Perlin and D. Fox. Pad: An Alternative Approach to the Computer Interface. Proc. of the 20th International Conference on COmputer Graphics and Interactive Techniques, August 1993, pp. 57-64.

[18] Sloan Digital Sky Survey. [online] Available: http://www.sdss.org/, cited February 2002.

[19] Spotfire Inc. [online] Available: http://www.spotfire.com, cited February 2002.

[20] C. Stolte, D. Tang, and P. Hanrahan. Polaris: A System for Query, Analysis, and Visualization of Multi-dimensional Relational Databases. In IEEE Transactions on Visualization and Computer Graphics, Vol. 8, No. 1, January 2002, pp. 52-65.

[21] D.Tang and M. Baker. Analysis of a Local-Area Wireless Network. In Proc. of the 6th International Conference on Mobile Computing and Networking, August 2000, pp. 1-10.

[22] K. Thearling, B. Becker, D. DeCoste, B. Mawby, M. Pilote, and D. Sommerfield. Visualizing Data Mining Models. In Information Visualization in Data Mining and Knowledge Discovery. Edited by U. Fayyad, G. Grinstein, and A. Wierse. Morgan Kaufman, 2001.

[23] E. Thomsen. OLAP Solutions: Building Multidimensional Information Systems. Wiley Computer Publishing, New York, 1997.

[24] M. O. Ward. XmdvTool: Integrating multiple methods for visualizing multi-variate data. In Proceedings of IEEE Visualization, 1994, pp. 326-336.

[25] J. Welling and M. Derthick. Visualization of Large Multi-dimensional Datasets. In Proc. of Virtual Observatories of the Future 2000.

[26] A. Woodruff, C. Olston, A. Aiken, M. Chu, V. Ercegovac, M. Lin, M. Spalding and M. Stonebraker. DataSplash: A Direct Manipulation Environment for Programming Semantic Zoom Visualizations of Tabular Data. Journal of Visual Languages and Computing, Special Issue on Visual Languages for End-user and Domain-specific Programming, 12(5), October 2001, pp. 551-571.